

kernel 3.18.7

in arch/arm/kernel/entry-armv.S

该文件中包括的是ARM core产生excption后的entries.即无论是interrupt还是page fault等等，当发生了，软件的入口就都在该文件中。

ARM instruction相比intel instruction更难读，而且gnu assembler的汇编格式相比Microsoft汇编格式也更难读(最起码有点i386汇编阅读史的我是这么认为的)。要看懂

entry-arm.S中的汇编还不如看vmlinux的反汇编码更容易。

```
$ arm-linux-gnueabi-objdump -d vmlinux-3.18.7-yocto-standard > linux-3.18.7.src
```

Disassembly of section .vectors:

下面就是entry-armv.S中的汇编源码被编译成code以后的反汇编code，我看着比源码要容易理解多了。那些gnu assembler的伪码看着就让人头疼。我以前写了2年的Microsoft

格式的x86汇编，从没有对汇编产生过太大的讨厌，但我每当看到gnu assembler格式的汇编就有种厌恶感，再加上是ARM instuction，更是一看就有想吐的感觉。或许是先入为

主的缘故，但一直不理解gnu那帮人为什么设计出这种format。

=====

Disassembly of section .vectors:

00000000 <__vectors_start>:

0: ea0003ff b 1004 <vector_rst>

```

4: ea000465    b    11a0 <vector_und>

8: e59ffff0  ldr  pc, [pc, #4080]    ; 1000 <__stubs_start>

c: ea000443    b    1120 <vector_pabt>

10: ea000422    b    10a0 <vector_dabt>

14: ea000481    b    1220 <vector_addrxcptn>

18: ea000400    b    1020 <vector_irq>

1c: ea000487    b    1240 <vector_fiq_offset>

```

Disassembly of section .stubs:

00001000 <__stubs_start>:

```

1000: c000e760    .word    0xc000e760

```

00001004 <vector_rst>:

```

1004: ef9f0000  svc  0x009f0000

1008: ea000064    b    11a0 <vector_und>

100c: e320f000    nop  {0}

1010: e320f000    nop  {0}

1014: e320f000    nop  {0}

1018: e320f000    nop  {0}

101c: e320f000    nop  {0}

```

00001020 <vector_irq>:

| | | | |
|-------|----------|-------|----------------------|
| 1020: | e24ee004 | sub | lr, lr, #4 |
| 1024: | e88d4001 | stm | sp, {r0, lr} |
| 1028: | e14fe000 | mrs | lr, SPSR |
| 102c: | e58de008 | str | lr, [sp, #8] |
| 1030: | e10f0000 | mrs | r0, CPSR |
| 1034: | e2200001 | eor | r0, r0, #1 |
| 1038: | e16ff000 | msr | SPSR_fsxc, r0 |
| 103c: | e20ee00f | and | lr, lr, #15 |
| 1040: | e1a0000d | mov | r0, sp |
| 1044: | e79fe10e | ldr | lr, [pc, lr, lsl #2] |
| 1048: | e1b0f00e | movs | pc, lr |
| 104c: | c00122c0 | .word | 0xc00122c0 |
| 1050: | c0011f60 | .word | 0xc0011f60 |
| 1054: | c0011f60 | .word | 0xc0011f60 |
| 1058: | c0012000 | .word | 0xc0012000 |
| 105c: | c0011f60 | .word | 0xc0011f60 |
| 1060: | c0011f60 | .word | 0xc0011f60 |
| 1064: | c0011f60 | .word | 0xc0011f60 |
| 1068: | c0011f60 | .word | 0xc0011f60 |
| 106c: | c0011f60 | .word | 0xc0011f60 |
| 1070: | c0011f60 | .word | 0xc0011f60 |
| 1074: | c0011f60 | .word | 0xc0011f60 |
| 1078: | c0011f60 | .word | 0xc0011f60 |
| 107c: | c0011f60 | .word | 0xc0011f60 |
| 1080: | c0011f60 | .word | 0xc0011f60 |

| | | | |
|-------|----------|-------|------------|
| 1084: | c0011f60 | .word | 0xc0011f60 |
| 1088: | c0011f60 | .word | 0xc0011f60 |
| 108c: | e320f000 | nop | {0} |
| 1090: | e320f000 | nop | {0} |
| 1094: | e320f000 | nop | {0} |
| 1098: | e320f000 | nop | {0} |
| 109c: | e320f000 | nop | {0} |

000010a0 <vector_dabt>:

| | | | |
|-------|----------|-------|----------------------|
| 10a0: | e24ee008 | sub | lr, lr, #8 |
| 10a4: | e88d4001 | stm | sp, {r0, lr} |
| 10a8: | e14fe000 | mrs | lr, SPSR |
| 10ac: | e58de008 | str | lr, [sp, #8] |
| 10b0: | e10f0000 | mrs | r0, CPSR |
| 10b4: | e2200004 | eor | r0, r0, #4 |
| 10b8: | e16ff000 | msr | SPSR_fsxc, r0 |
| 10bc: | e20ee00f | and | lr, lr, #15 |
| 10c0: | e1a0000d | mov | r0, sp |
| 10c4: | e79fe10e | ldr | lr, [pc, lr, lsl #2] |
| 10c8: | e1b0f00e | movs | pc, lr |
| 10cc: | c0012280 | .word | 0xc0012280 |
| 10d0: | c0011f50 | .word | 0xc0011f50 |
| 10d4: | c0011f50 | .word | 0xc0011f50 |
| 10d8: | c0011fa0 | .word | 0xc0011fa0 |
| 10dc: | c0011f50 | .word | 0xc0011f50 |

| | | | |
|-------|----------|-------|------------|
| 10e0: | c0011f50 | .word | 0xc0011f50 |
| 10e4: | c0011f50 | .word | 0xc0011f50 |
| 10e8: | c0011f50 | .word | 0xc0011f50 |
| 10ec: | c0011f50 | .word | 0xc0011f50 |
| 10f0: | c0011f50 | .word | 0xc0011f50 |
| 10f4: | c0011f50 | .word | 0xc0011f50 |
| 10f8: | c0011f50 | .word | 0xc0011f50 |
| 10fc: | c0011f50 | .word | 0xc0011f50 |
| 1100: | c0011f50 | .word | 0xc0011f50 |
| 1104: | c0011f50 | .word | 0xc0011f50 |
| 1108: | c0011f50 | .word | 0xc0011f50 |
| 110c: | e320f000 | nop | {0} |
| 1110: | e320f000 | nop | {0} |
| 1114: | e320f000 | nop | {0} |
| 1118: | e320f000 | nop | {0} |
| 111c: | e320f000 | nop | {0} |

00001120 <vector_pabt>:

| | | | |
|-------|----------|-----|---------------|
| 1120: | e24ee004 | sub | lr, lr, #4 |
| 1124: | e88d4001 | stm | sp, {r0, lr} |
| 1128: | e14fe000 | mrs | lr, SPSR |
| 112c: | e58de008 | str | lr, [sp, #8] |
| 1130: | e10f0000 | mrs | r0, CPSR |
| 1134: | e2200004 | eor | r0, r0, #4 |
| 1138: | e16ff000 | msr | SPSR_fsxc, r0 |

| | | |
|-------|----------|--------------------------|
| 113c: | e20ee00f | and lr, lr, #15 |
| 1140: | e1a0000d | mov r0, sp |
| 1144: | e79fe10e | ldr lr, [pc, lr, lsl #2] |
| 1148: | e1b0f00e | movs pc, lr |
| 114c: | c00124e0 | .word 0xc00124e0 |
| 1150: | c0011f40 | .word 0xc0011f40 |
| 1154: | c0011f40 | .word 0xc0011f40 |
| 1158: | c0012120 | .word 0xc0012120 |
| 115c: | c0011f40 | .word 0xc0011f40 |
| 1160: | c0011f40 | .word 0xc0011f40 |
| 1164: | c0011f40 | .word 0xc0011f40 |
| 1168: | c0011f40 | .word 0xc0011f40 |
| 116c: | c0011f40 | .word 0xc0011f40 |
| 1170: | c0011f40 | .word 0xc0011f40 |
| 1174: | c0011f40 | .word 0xc0011f40 |
| 1178: | c0011f40 | .word 0xc0011f40 |
| 117c: | c0011f40 | .word 0xc0011f40 |
| 1180: | c0011f40 | .word 0xc0011f40 |
| 1184: | c0011f40 | .word 0xc0011f40 |
| 1188: | c0011f40 | .word 0xc0011f40 |
| 118c: | e320f000 | nop {0} |
| 1190: | e320f000 | nop {0} |
| 1194: | e320f000 | nop {0} |
| 1198: | e320f000 | nop {0} |
| 119c: | e320f000 | nop {0} |

000011a0 <vector_und>:

| | | |
|-------|----------|--------------------------|
| 11a0: | e88d4001 | stm sp, {r0, lr} |
| 11a4: | e14fe000 | mrs lr, SPSR |
| 11a8: | e58de008 | str lr, [sp, #8] |
| 11ac: | e10f0000 | mrs r0, CPSR |
| 11b0: | e2200008 | eor r0, r0, #8 |
| 11b4: | e16ff000 | msr SPSR_fsrc, r0 |
| 11b8: | e20ee00f | and lr, lr, #15 |
| 11bc: | e1a0000d | mov r0, sp |
| 11c0: | e79fe10e | ldr lr, [pc, lr, lsl #2] |
| 11c4: | e1b0f00e | movs pc, lr |
| 11c8: | c0012320 | .word 0xc0012320 |
| 11cc: | c0011f70 | .word 0xc0011f70 |
| 11d0: | c0011f70 | .word 0xc0011f70 |
| 11d4: | c00120a0 | .word 0xc00120a0 |
| 11d8: | c0011f70 | .word 0xc0011f70 |
| 11dc: | c0011f70 | .word 0xc0011f70 |
| 11e0: | c0011f70 | .word 0xc0011f70 |
| 11e4: | c0011f70 | .word 0xc0011f70 |
| 11e8: | c0011f70 | .word 0xc0011f70 |
| 11ec: | c0011f70 | .word 0xc0011f70 |
| 11f0: | c0011f70 | .word 0xc0011f70 |
| 11f4: | c0011f70 | .word 0xc0011f70 |
| 11f8: | c0011f70 | .word 0xc0011f70 |

| | | | |
|-------|----------|-------|------------|
| 11fc: | c0011f70 | .word | 0xc0011f70 |
| 1200: | c0011f70 | .word | 0xc0011f70 |
| 1204: | c0011f70 | .word | 0xc0011f70 |
| 1208: | e320f000 | nop | {0} |
| 120c: | e320f000 | nop | {0} |
| 1210: | e320f000 | nop | {0} |
| 1214: | e320f000 | nop | {0} |
| 1218: | e320f000 | nop | {0} |
| 121c: | e320f000 | nop | {0} |

00001220 <vector_addrxcptn>:

| | | | |
|-------|-----------|-----|-------------------------|
| 1220: | eaffffffe | b | 1220 <vector_addrxcptn> |
| 1224: | e320f000 | nop | {0} |
| 1228: | e320f000 | nop | {0} |
| 122c: | e320f000 | nop | {0} |
| 1230: | e320f000 | nop | {0} |
| 1234: | e320f000 | nop | {0} |
| 1238: | e320f000 | nop | {0} |
| 123c: | e320f000 | nop | {0} |

00001240 <vector_fiq_offset>:

| | | | |
|-------|----------|-----|--------------|
| 1240: | e24ee004 | sub | lr, lr, #4 |
| 1244: | e88d4001 | stm | sp, {r0, lr} |
| 1248: | e14fe000 | mrs | lr, SPSR |
| 124c: | e58de008 | str | lr, [sp, #8] |

| | | |
|-------|----------|--------------------------|
| 1250: | e10f0000 | mrs r0, CPSR |
| 1254: | e2200002 | eor r0, r0, #2 |
| 1258: | e16ff000 | msr SPSR_fsrc, r0 |
| 125c: | e20ee00f | and lr, lr, #15 |
| 1260: | e1a0000d | mov r0, sp |
| 1264: | e79fe10e | ldr lr, [pc, lr, lsl #2] |
| 1268: | e1b0f00e | movs pc, lr |
| 126c: | c0012540 | .word 0xc0012540 |
| 1270: | c0012180 | .word 0xc0012180 |
| 1274: | c0012180 | .word 0xc0012180 |
| 1278: | c0012180 | .word 0xc0012180 |
| 127c: | c0012180 | .word 0xc0012180 |
| 1280: | c0012180 | .word 0xc0012180 |
| 1284: | c0012180 | .word 0xc0012180 |
| 1288: | c0012200 | .word 0xc0012200 |
| 128c: | c0012180 | .word 0xc0012180 |
| 1290: | c0012180 | .word 0xc0012180 |
| 1294: | c0012180 | .word 0xc0012180 |
| 1298: | c0012180 | .word 0xc0012180 |
| 129c: | c0012180 | .word 0xc0012180 |
| 12a0: | c0012180 | .word 0xc0012180 |
| 12a4: | c0012180 | .word 0xc0012180 |
| 12a8: | c0012180 | .word 0xc0012180 |
| 12ac: | e320f000 | nop {0} |
| 12b0: | e320f000 | nop {0} |

12b4: e320f000 nop {0}

12b8: e320f000 nop {0}

12bc: e320f000 nop {0}

=====

上面的code分别位于".vectors"和".stubs" section中。这个在entry-armv.S源码中反映了。

```
.section .vectors, "ax", %progbits    ①
```

```
__vectors_start:
```

```
W(b)    vector_rst
```

```
W(b)    vector_und
```

```
W(ldr)   pc, __vectors_start + 0x1000
```

```
W(b)    vector_pabt
```

```
W(b)    vector_dabt
```

```
W(b)    vector_addrexcptn
```

```
W(b)    vector_irq
```

```
W(b)    vector_fiq
```

这是".vector" section

```
.section .stubs, "ax", %progbits    ②
```

```
__stubs_start:
```

```
@ This must be the first word
```

```
.word    vector_swi
```

vector_rst:

ARM(swi SYS_ERROR0)

THUMB(svc #0)

THUMB(nop)

b vector_und

/*

* Interrupt dispatcher

*/

vector_stub irq, IRQ_MODE, 4

.long__irq_usr @ 0 (USR_26 / USR_32)

.long__irq_invalid @ 1 (FIQ_26 / FIQ_32)

.long__irq_invalid @ 2 (IRQ_26 / IRQ_32)

.long__irq_svc @ 3 (SVC_26 / SVC_32)

.long__irq_invalid @ 4

.long__irq_invalid @ 5

.long__irq_invalid @ 6

.long__irq_invalid @ 7

.long__irq_invalid @ 8

.long__irq_invalid @ 9

.long__irq_invalid @ a

.long__irq_invalid @ b

.long__irq_invalid @ c

.long__irq_invalid @ d

.long__irq_invalid @ e

.long__irq_invalid @ f

/*

* Data abort dispatcher

* Enter in ABT mode, spsr = USR CPSR, lr = USR PC

*/

vector_stub dabt, ABT_MODE, 8

.....

①

.vector是section name,"ax" 表示是relocatable and executable

②

.stubs是section name,"ax" 表示是relocatable and executable

".vectors" section被安排在0地址 (page 0 in virtual space) , 而".stubs" section则被安排在0x1000(page 1 in virtual space).

从反汇编码看, 这一目了然, 而从源码上是看不出的, 要从vmlinux.lds的连接脚本上才能看出。

in vmlinux.lds

/*

* The vectors and stubs are relocatable code, and the

* only thing that matters is their relative offsets

*/

```
__vectors_start = .;
```

```
.vectors 0 : AT(__vectors_start) {
```

```
*(.vectors)
```

```
}
```

```
. = __vectors_start + SIZEOF(.vectors);
```

```
__vectors_end = .;
```

```
__stubs_start = .;
```

```
.stubs 0x1000 : AT(__stubs_start) {
```

```
*(.stubs)
```

```
}
```

```
. = __stubs_start + SIZEOF(.stubs);
```

```
__stubs_end = .;
```

从build kernel后生成的system.map文件也能看出端倪。

in config-3.18.7-yocto-standard

```
00000000 t __vectors_start
```

```
00000024 A cpu_v7_suspend_size
```

```
0000002c A cpu_ca9mp_suspend_size
```

```
00001000 t __stubs_start
```

```
00001004 t vector_rst
```

00001020 t vector_irq

000010a0 t vector_dabt

00001120 t vector_pabt

000011a0 t vector_und

00001220 t vector_addrxcptn

00001240 t vector_fiq

00001240 T vector_fiq_offset

位于0地址的".vectors" section只是包含了ARM core的中断向量表。依次为

1. RESET
2. UNDEF (undefined instruction, a.k.a. invalid instruction)
3. SWI (SoftWare Interrupt)
4. PABT (Prefetch instruction Abort)
5. DABT (Data Abort)
6. - ?
7. IRQ hardware interrupt
8. FIQ Fast Interrupt

第6项在“ARM System Developer's Guide: Designing and Optimizing System Software” book中被标为Reserved。这本书只讲了RAMv5，在ARMv7-A manual中标为：

0x14 Not used

其中SWI (SoftWare Interrupt)有点特殊。

```
W(ldr)    pc, __vectors_start + 0x1000
```

反汇编后为

```
ldr    pc, [pc, #4080]    ; 1000 <__stubs_start>
```

反汇编码比源码更易理解。

pc, #4080 指向__stubs_start的开始。

__stubs_start:

@ This must be the first word

```
.word    vector_swi
```

__stubs_start的头四个bytes是vector_swi的地址，所以

ldr pc, [pc, #4080] 就是把vector_swi赋值给pc，也就是跳转到vector_swi执行。

位于".stubs" section的都是第一层exception handler。

".vectors" section所在的page实际上就0x20(32)个bytes是有意义的，而page 1所在的".stubs" section包含的first level exception handlers code不过0x2e0

in config-3.18.7-yocto-standard

c05a9000 T __vectors_start

c05a9020 T __stubs_start

c05a9020 T __vectors_end

c05a92e0 t __mmap_switched

c05a92e0 T __stubs_end

0xc05a92e0 - 0xc05a9020 = 0x2e0

1. exception vector table

| | | | | |
|-----|----------|-----|-----------------|------------------------|
| 0: | ea0003ff | b | 1004 | <vector_rst> |
| 4: | ea000465 | b | 11a0 | <vector_und> |
| 8: | e59ffff0 | ldr | pc, [pc, #4080] | ; 1000 <__stubs_start> |
| c: | ea000443 | b | 1120 | <vector_pabt> |
| 10: | ea000422 | b | 10a0 | <vector_dabt> |
| 14: | ea000481 | b | 1220 | <vector_addrxcptn> |
| 18: | ea000400 | b | 1020 | <vector_irq> |
| 1c: | ea000487 | b | 1240 | <vector_fiq_offset> |

当出现exception，ARM core就会载入对应的entry执行。比如reset ARM core,core将跳转到0地址(当vector在0地址的情况下)，也就是这里执行

b 1004 <vector_rst>

2. vector stub code

比如当hardware interrupt发生时，core将执行

18: ea000400 b 1020 <vector_irq>

跳转到下面的位于vector stub code中的irq entry

00001020 <vector_irq>:

| | | | |
|-------|----------|--------------------------|---|
| 1020: | e24ee004 | sub lr, lr, #4 | ① |
| 1024: | e88d4001 | stm sp, {r0, lr} | ② |
| 1028: | e14fe000 | mrs lr, SPSR | |
| 102c: | e58de008 | str lr, [sp, #8] | |
| 1030: | e10f0000 | mrs r0, CPSR | ③ |
| 1034: | e2200001 | eor r0, r0, #1 | |
| 1038: | e16ff000 | msr SPSR_fsrc, r0 | |
| 103c: | e20ee00f | and lr, lr, #15 | |
| 1040: | e1a0000d | mov r0, sp | ④ |
| 1044: | e79fe10e | ldr lr, [pc, lr, lsl #2] | |
| 1048: | e1b0f00e | movs pc, lr | |

| | | | |
|-------|----------|-------|------------|
| 104c: | c00122c0 | .word | 0xc00122c0 |
| 1050: | c0011f60 | .word | 0xc0011f60 |
| 1054: | c0011f60 | .word | 0xc0011f60 |
| 1058: | c0012000 | .word | 0xc0012000 |
| 105c: | c0011f60 | .word | 0xc0011f60 |
| 1060: | c0011f60 | .word | 0xc0011f60 |
| 1064: | c0011f60 | .word | 0xc0011f60 |
| 1068: | c0011f60 | .word | 0xc0011f60 |
| 106c: | c0011f60 | .word | 0xc0011f60 |
| 1070: | c0011f60 | .word | 0xc0011f60 |
| 1074: | c0011f60 | .word | 0xc0011f60 |
| 1078: | c0011f60 | .word | 0xc0011f60 |
| 107c: | c0011f60 | .word | 0xc0011f60 |
| 1080: | c0011f60 | .word | 0xc0011f60 |
| 1084: | c0011f60 | .word | 0xc0011f60 |
| 1088: | c0011f60 | .word | 0xc0011f60 |

上面的反汇编码对应的是下面的源码

```
/*  
  
* Interrupt dispatcher  
  
*/  
  
vector_stub    irq, IRQ_MODE, 4
```

| | |
|--------------------|-----------------------|
| .long__irq_usr | @ 0 (USR_26 / USR_32) |
| .long__irq_invalid | @ 1 (FIQ_26 / FIQ_32) |
| .long__irq_invalid | @ 2 (IRQ_26 / IRQ_32) |
| .long__irq_svc | @ 3 (SVC_26 / SVC_32) |
| .long__irq_invalid | @ 4 |
| .long__irq_invalid | @ 5 |
| .long__irq_invalid | @ 6 |
| .long__irq_invalid | @ 7 |
| .long__irq_invalid | @ 8 |
| .long__irq_invalid | @ 9 |
| .long__irq_invalid | @ a |
| .long__irq_invalid | @ b |
| .long__irq_invalid | @ c |
| .long__irq_invalid | @ d |
| .long__irq_invalid | @ e |
| .long__irq_invalid | @ f |

vector_stub macro下面是不同mode下的excetion handler，这里就是不同mode下的interrupt handler。

由于Linux只工作在user mode和svc mode，所以上面只有user mode与svc mode对应的handler才assign有意义的handler。

ARM只定义了7种处理器mode，由core的cpsr register的最低5位(bit 0 to bit 4)指示CPU mode

in arch/arm/include/uapi/asm/ptrace.h

```
#define USR_MODE      0x00000010
#define SVC_MODE      0x00000013

#define FIQ_MODE      0x00000011
#define IRQ_MODE      0x00000012
#define ABT_MODE      0x00000017
#define HYP_MODE      0x0000001a
#define UND_MODE      0x0000001b
#define SYSTEM_MODE    0x0000001f
```

由于指示mode的bit 4总是1（从上面的定义可知），所以实际上用bit 0 to bit 3可以标示which mode。

上面的mode exception handler共定义了16项（所有mode全用上也不过7项，何况Linux只使用了2项），大部分都填的是__irq_invalid。

user mode和svc mode的在cpsr register中的mode bit pattern是0x10和0x13,去掉bit 4的“1”,值是0x00和0x03,他们正好是mode exception handler table中的

第0项和第3项。

①

```
vector_stub    irq, IRQ_MODE, 4
```

这里的4是对lr regsiter的调整。由于RISC CPU的特性，pc register总是先于executing的instruction. 在ARM上是先于2条指令，即8 bytes的offset。每条指令4个bytes.

instruction 1

instruction 2 <--- executing , interrupt occur

instruction 3

instruction 4 <--- pc 指向

instruction 5

instruction 6

由于是在执行instruction 2时发生interrupt,当core切换入interrupt handler时，instruction 2必须执行完(hardware interrupt本身是不能中断正在执行的instruction的，它只能在本指令执行完后再打断)，所以instruction 3是在interrupt handler完成以后应该执行的instruction，即hardware interrupt handler在instruction 2和instruction 3间执行。lr用于记录interrupt handler处理完后要回复到原执行流程的地址。而instruction 3的地址正好是pc - 4。这就是这里4的由来。

vector_stub dabt, ABT_MODE, 8

这里的8则是对产生data abort exception后对lr的调整。它调整不同的值与interrupt是由于下面这样的特性。

interrupt是在执行完当前正在executing的instruction后进入interrupt handler，即当前executing instruction的效果已经结束，当interrupt handler结束，就执行下一条instruction，不能再执行上面那条指令了，否则就错了。而data abort exception之所以产生，就是因为当前executing的instruction没法运行下去（当然就不存在该指令的执行效果的问题），所以当exception handler返回，必须再次执行原来产生exception的指令。

instruction 1

instruction 2 <--- executing, data abort exception occur

instruction 3

instruction 4 <--- pc 指向

instruction 5

instruction 6

假设core在执行instruction 2时发生exception,比如该page不可写。则进入exception handler后, kernel的责任是fix该exception, 使得exception handler返回后依然

执行instruction 2.而instruction 2的地址就是pc - 8。

从上可知interrupt是异步的,而exception是同步的。这个在intel CPU manual中有分别,但在ARM manual中好像都统称为"exception"。

②

保存ro,lr和spsr registers to stack

③

disable IRQ and enter svc mode

④

根据发生exception发生时的mode跳转到对应的handler。

比如如果发生interrupt时,core运行在user mode,则跳转到__irq_usr handler,如果是svc mode,则跳转到__irq_svc。

由于Linux只工作在这两个mode,所以其它的任意mode,都作出错处理(__irq_invalid).

当kernel起来后,真正的exception vector table和first level exception handler并不在0地址开始,而是在0xfffff000。

in in config-3.18.7-yocto-standard

CONFIG_VECTORS_BASE=0xffff0000

整个move exception vector table和first level exception handler的代码在early_trap_init() in arch/arm/kernel/traps.c.

start_kernel()

|

|

\\

setup_arch() in arch/arm/kernel/setup.c

|

|

\\

paging_init() in arch/arm/mm/mmu.c

|

|

\\

devicemaps_init() in arch/arm/mm/mmu.c

|

|

\\

early_trap_init() in arch/arm/mm/mmu.c

in devicemaps_init()

```
/*  
 * Allocate the vector page early.  
 */  
vectors = early_alloc(PAGE_SIZE * 2); ①  
  
early_trap_init(vectors);
```

①

由于kernel的memory manager(比如buddy allocator, slab allocator, etc)都没初始化, 这里分配2 pages的memory用的是memblock module(当kernel起来后就丢弃不用了)

```
void __init early_trap_init(void *vectors_base)  
{  
#ifndef CONFIG_CPU_V7M ②  
    unsigned long vectors = (unsigned long)vectors_base;  
    extern char __stubs_start[], __stubs_end[]; ③  
    extern char __vectors_start[], __vectors_end[];  
    unsigned i;  
  
    vectors_page = vectors_base;  
  
    /*
```


- * Poison the vectors page with an undefined instruction. This
- * instruction is chosen to be undefined for both ARM and Thumb
- * ISAs. The Thumb version is an undefined instruction with a
- * branch back to the undefined instruction.

*/

for (i = 0; i < PAGE_SIZE / sizeof(u32); i++) ④

((u32 *)vectors_base)[i] = 0xe7fdddef1;

/*

- * Copy the vectors, stubs and kuser helpers (in entry-armv.S)
- * into the vector page, mapped at 0xffff0000, and ensure these
- * are visible to the instruction stream.

*/

memcpy((void *)vectors, __vectors_start, __vectors_end - __vectors_start); ⑤

memcpy((void *)vectors + 0x1000, __stubs_start, __stubs_end - __stubs_start); ⑥

kuser_init(vectors_base); ⑦

flush_icache_range(vectors, vectors + PAGE_SIZE * 2); ⑧

modify_domain(DOMAIN_USER, DOMAIN_CLIENT);

#else /* ifndef CONFIG_CPU_V7M */

/*

- * on V7-M there is no need to copy the vector table to a dedicated
- * memory area. The address is configurable and so a table in the kernel
- * image can be used.

```
*/  
  
#endif  
  
}
```

②

CONFIG_CPU_V7M is for Cortex-M, 当然没有定义该macro。

③

这些extern的global variable都定义在vmlinux.lds中

```
. = __vectors_start + SIZEOF(.vectors);  
  
__vectors_end = .;  
  
__stubs_start = .;  
  
.stubs 0x1000 : AT(__stubs_start) {  
  
    *(.stubs)  
  
}  
  
. = __stubs_start + SIZEOF(.stubs);  
  
__stubs_end = .;
```

④

从comments可知，是把存放exception vector table的page先填满invalid instruction.

⑤

把exception vector table复制过去，这样除了头上32个字节是exception vector table，该page的其他内容都是invalid instruction。

⑥

把first level exception handler的代码复制到第2页.

⑦

CONFIG_KUSER_HELPERS=y

in arch/arm/kernel/traps.c

```
#ifdef CONFIG_KUSER_HELPERS
```

```
static void __init kuser_init(void *vectors)
```

```
{
```

```
    extern char __kuser_helper_start[], __kuser_helper_end[];
```

```
    int kuser_sz = __kuser_helper_end - __kuser_helper_start;
```

```
    memcpy(vectors + 0x1000 - kuser_sz, __kuser_helper_start, kuser_sz);
```

```
    /*
```

```
     * vectors + 0xfe0 = __kuser_get_tls
```

```
     * vectors + 0xfe8 = hardware TLS instruction at 0xffff0fe8
```

```
    */
```

```
    if (tls_emu || has_tls_reg)
```

```
        memcpy(vectors + 0xfe0, vectors + 0xfe8, 4);
```

```
}
```

```
#else
```

```
static inline void __init kuser_init(void *vectors)
```

```
{
```

```
}
```

```
#endif
```

__kuser_helper_start and __kuser_helper_end 定义在arch/arm/kernel/entry-armv.S

in system.map

```
c05acb80 t __kuser_cmpxchg64
```

```
c05acb80 T __kuser_helper_start
```

```
c05acbc0 t __kuser_memory_barrier
```

```
c05acbe0 t __kuser_cmpxchg
```

```
c05acc00 t __kuser_get_tls
```

```
c05acc1c t __kuser_helper_version
```

```
c05acc20 T __kuser_helper_end
```

```
c05acc20 - c05acb80 = 0xa0 (160 bytes)
```

kuser_init()把从__kuser_helper_start to __kuser_helper_end的代码复制到vector page的尾部。

这样vector page的头上是32个bytes的exception vectors(8 entries),尾部是160bytes的kuser helper code,而中间被填充了invalid instruction(0xe7fdddef1)

具体kuser helper的分析,见《about kuser helper》笔记。

由于无论是vector page中的code(exception table中的跳转指令或是kuser helper code),还是vector stub page中的code都生成在0地址开始的virtual address,

但都被复制到0xffff,f000开始的地址去运行, 所以这些code必须是relocatable code, 即这些code本身间的跳转不能有绝对地址的跳转, 只能有相对地址的跳转。当然对这些relocatable code本身外的跳转是完全可以使用绝对地址的。

mapping vector page to virtual address **0xffff,0000**

and vector stub page to **0xffff,1000**

```
static void __init devicemaps_init(const struct machine_desc *mdesc)
```

```
{
```

```
    struct map_desc map;
```

```
    unsigned long addr;
```

```
    void *vectors;
```

```
    /*
```

```
     * Allocate the vector page early.
```

```
     */
```

```
    vectors = early_alloc(PAGE_SIZE * 2);
```

```
    early_trap_init(vectors);
```

```
    for (addr = VMALLOC_START; addr; addr += PMD_SIZE)
```

```
        pmd_clear(pmd_off_k(addr));
```

```
    .....
```

```
/*
```

```
 * Create a mapping for the machine vectors at the high-vectors
```

```
 * location (0xffff0000). If we aren't using high-vectors, also
```

```
 * create a mapping at the low-vectors virtual address.
```

```
*/
```

```
map.pfn = __phys_to_pfn(virt_to_phys(vectors)); ①
```

```
map.virtual = 0xffff0000;
```

```
map.length = PAGE_SIZE;
```

```
#ifdef CONFIG_KUSER_HELPERS
```

```
    map.type = MT_HIGH_VECTORS;
```

```
#else
```

```
    map.type = MT_LOW_VECTORS;
```

```
#endif
```

```
    create_mapping(&map);
```

```
    if (!vectors_high()) { ②
```

```
        map.virtual = 0;
```

```
        map.length = PAGE_SIZE * 2;
```

```
        map.type = MT_LOW_VECTORS;
```

```
        create_mapping(&map);
```

```
    }
```

```
/* Now create a kernel read-only mapping */
```

```
map.pfn += 1; ③
```

```

map.virtual = 0xffff0000 + PAGE_SIZE;

map.length = PAGE_SIZE;

map.type = MT_LOW_VECTORS;

create_mapping(&map);

.....

}

```

①

mapping vector page to 0xffff0000,建立virtual page v.s. physical page之间的page table.

②

```

#if __LINUX_ARM_ARCH__ >= 4

#define vectors_high() (get_cr() & CR_V)

#else

#define vectors_high() (0)

#endif

#define CR_V (1 << 13) /* Vectors relocated to 0xffff0000 */

```

In ARMv7-A manual,

If the Security Extensions are not implemented there is a single exception base address. This is controlled by the SCTL.R.V bit:

V == 0 Exception base address = 0x00000000. This setting is referred to as normal vectors, or as low

vectors.

V == 1 Exception base address = 0xFFFF0000. This setting is referred to as high vectors, or Hivecs.

Note

Use of the Hivecs setting, V == 1, is deprecated in ARMv7-R. ARM recommends that Hivecs is used only

in ARMv7-A implementations.

```
static inline unsigned long get_cr(void)
```

```
{  
    unsigned long val;  
    asm("mrc p15, 0, %0, c1, c0, 0    @ get CR" : "=r" (val) : : "cc");  
    return val;  
}
```

mapping vector page and vector stub page 2 virtual address 0 and 0x1000

③

mapping vector stub page to 0xffff0000 + 0x1000