

当device / driver加入到bus上时，bus会发出对应的notification,kernel code可以关注这些notification，并对device或driver进行修改！

in linux/device.h

```
1.  /* All 4 notifiers below get called with the target struct device *
2.   * as an argument. Note that those functions are likely to be called
3.   * with the device lock held in the core, so be careful.
4.   */
5.  #define BUS_NOTIFY_ADD_DEVICE      0x00000001 /* device added */
6.  #define BUS_NOTIFY_DEL_DEVICE      0x00000002 /* device to be removed */
7.  #define BUS_NOTIFY_REMOVED_DEVICE  0x00000003 /* device removed */
8.  #define BUS_NOTIFY_BIND_DRIVER     0x00000004 /* driver about to be
9.                                           bound */
10. #define BUS_NOTIFY_BOUND_DRIVER     0x00000005 /* driver bound to device */
11. #define BUS_NOTIFY_UNBIND_DRIVER    0x00000006 /* driver about to be
12.                                           unbound */
13. #define BUS_NOTIFY_UNBOUND_DRIVER   0x00000007 /* driver is unbound
14.                                           from the device */
```

注册关注这些notification的APIs

```
1.  /*
2.   * Bus notifiers: Get notified of addition/removal of devices
3.   * and binding/unbinding of drivers to devices.
4.   * In the long run, it should be a replacement for the platform
5.   * notify hooks.
6.   */
7.  struct notifier_block;
8.
9.  extern int bus_register_notifier(struct bus_type *bus,
10.                                struct notifier_block *nb);
11.  extern int bus_unregister_notifier(struct bus_type *bus,
12.                                   struct notifier_block *nb);
```

kernel 4.2.8中pegmatite SoC的 `drivers/platform/pegmatite/dma.c` 就利用了这种机制来实现对dma memory分配的physical address是否大于2G进行检查。

```

1. #include <linux/kernel.h>
2. #include <linux/init.h>
3. #include <linux/io.h>
4. #include <linux/dma-mapping.h>
5. #include <linux/sizes.h>
6. #include <linux/pci.h>
7. #include <linux/platform_device.h>
8.
9. static struct dma_map_ops pegmatite_dma_ops;
10. static struct dma_map_ops *orig_dma_ops;
11. static u64 default_dmamask = DMA_BIT_MASK(31);
12.
13. struct platform_device dummy_default_dma_dev = {
14.     .name    = "dummy_default_dma",
15.     .id     = -1,
16.     .dev     = {
17.         .dma_mask      = &default_dmamask,
18.         .coherent_dma_mask = DMA_BIT_MASK(31),
19.     },
20. };
21.
22. static void *pegmatite_alloc(struct device *dev, size_t size,
23.                             dma_addr_t *dma_handle, gfp_t gfp,
24.                             struct dma_attrs *attrs)
25. {
26.     if (WARN_ON_ONCE(!dev)) {
27.         dev = &dummy_default_dma_dev.dev;
28.     }
29.     return orig_dma_ops->alloc(dev, size, dma_handle, gfp, attrs);
30. }
31.
32. static void pegmatite_free(struct device *dev, size_t size,
33.                            void *vaddr, dma_addr_t dma_handle,
34.                            struct dma_attrs *attrs)
35. {
36.     if (WARN_ON_ONCE(!dev)) {
37.         dev = &dummy_default_dma_dev.dev;
38.     }
39.     orig_dma_ops->free(dev, size, vaddr, dma_handle, attrs);
40. }
41.
42. /* These DMA functions should behave the same as the generic ARM
43.  * functions, but with an additional sanity check to verify that we
44.  * don't try to DMA to DRAM above 2GB, because HW can't address that
45.  * memory. */
46. static dma_addr_t pegmatite_map_page(struct device *dev, struct page *page,
47.                                       unsigned long offset, size_t size,
48.                                       enum dma_data_direction dir,
49.                                       struct dma_attrs *attrs)
50. {
51.     BUG_ON(page_to_phys(page) >= SZ_2G);    ⑥
52.     if (WARN_ON_ONCE(!dev)) {
53.         dev = &dummy_default_dma_dev.dev;

```

```

54.     }
55.     return orig_dma_ops->map_page(dev, page, offset, size, dir, attrs);
56. }
57.
58. static int pegmatite_map_sg(struct device *dev, struct scatterlist *sg, int
nents,
59.                             enum dma_data_direction dir, struct dma_attrs *attrs)
60. {
61.     int i;
62.     struct scatterlist *s;
63.
64.     for_each_sg(sg, s, nents, i) {
65.         BUG_ON(page_to_phys(sg_page(s)) >= SZ_2G);
66.     }
67.
68.     if (WARN_ON_ONCE(!dev)) {
69.         dev = &dummy_default_dma_dev.dev;
70.     }
71.     return orig_dma_ops->map_sg(dev, sg, nents, dir, attrs);
72. }
73.
74. static void init_dma_ops (void)
75. {
76.     int rc;
77.
78.     orig_dma_ops = get_dma_ops(NULL);
79.     memcpy(&pegmatite_dma_ops, orig_dma_ops, sizeof(pegmatite_dma_ops));
80.     pegmatite_dma_ops.alloc = pegmatite_alloc;
81.     pegmatite_dma_ops.free = pegmatite_free;
82.     pegmatite_dma_ops.map_page = pegmatite_map_page;
83.     pegmatite_dma_ops.map_sg = pegmatite_map_sg;
84. #ifdef CONFIG_ARM64
85.     /* Set the global default dma ops */
86.     dma_ops = &pegmatite_dma_ops;
87. #endif
88.     rc = platform_device_register(&dummy_default_dma_dev);
89.     if (rc) {
90.         pr_err("Registering dma device failed! %d\n", rc);
91.         return;
92.     }
93. }
94.
95. static int dma_notifier(struct notifier_block *nb,
96.                         unsigned long event, void *__dev)
97. {
98.     struct device *dev = __dev;
99.
100.    if (event != BUS_NOTIFY_BIND_DRIVER) ③
101.        return NOTIFY_DONE;
102.    dev->archdata.dma_ops = &pegmatite_dma_ops; ④
103.
104.    if (strstr(dev_name(dev), "d0700000.stmmac")) { ⑤
105.        if (dev->dma_mask)
106.            *dev->dma_mask = DMA_BIT_MASK(31);

```

```

107.         dev->coherent_dma_mask = DMA_BIT_MASK(31);
108.     }
109.
110.     return NOTIFY_OK;
111. }
112.
113. static struct notifier_block dma_nb = {
114.     .notifier_call = dma_notifier,
115. };
116.
117. static int __init dma_init(void)
118. {
119.     init_dma_ops();
120.     bus_register_notifier(&platform_bus_type,    ①
121.                          &dma_nb);
122.     #if IS_ENABLED(CONFIG_PCI)
123.         bus_register_notifier(&pci_bus_type,      ②
124.                              &dma_nb);
125.     #endif
126.     return 0;
127. }
128. arch_initcall(dma_init);

```

①②

关注 `platform bus`和 `pci bus`上的driver。

③

只关心 `BUS_NOTIFY_BIND_DRIVER` notification , 即当driver与对应的device将要bind(绑定)之时。

④

替换原来的 `struct dma_map_ops`

⑤

可以针对特定的device

⑥

在map page以前做针对pegmatite SoC的检查。