

Virtual kernel memory layout:

vector : 0xffff0000 - 0xffff1000 (4 kB)

fixmap : 0xffc00000 - 0xffe00000 (2048 kB)

vmalloc : 0xf0000000 - 0xff000000 (240 MB)

lowmem : 0xc0000000 - 0xef800000 (760 MB)

pkmap : 0xbfe00000 - 0xc0000000 (2 MB)

modules : 0xbf000000 - 0xbfe00000 (14 MB)

.text : 0xc0008000 - 0xc05cde80 (5912 kB)

.init : 0xc05ce000 - 0xc0602000 (208 kB)

.data : 0xc0602000 - 0xc0638728 (218 kB)

.bss : 0xc0638728 - 0xc06a8af4 (449 kB)

在kernel初始化阶段

start_kernel() --> setup_arch() --> paging_init() --> devicemaps_init()

```
1. static void __init devicemaps_init(const struct machine_desc *mdesc)
2. {
3.     struct map_desc map;
4.     unsigned long addr;
5.     void *vectors;
6.
7.     /*
8.      * Allocate the vector page early.
9.      */
10.    vectors = early_alloc(PAGE_SIZE * 2);
11.
12.    early_trap_init(vectors);
13.
14.    for (addr = VMALLOC_START; addr; addr += PMD_SIZE)
15.        pmd_clear(pmd_off_k(addr));
16.
17.
18.    .....
19.
20.    /*
21.     * Create a mapping for the machine vectors at the high-vectors
22.     * location (0xffff0000). If we aren't using high-vectors, also
23.     * create a mapping at the low-vectors virtual address.
24.     */
25.    map.pfn = __phys_to_pfn(virt_to_phys(vectors));
26.    map.virtual = 0xffff0000;
27.    map.length = PAGE_SIZE;
28.    #ifdef CONFIG_KUSER_HELPERS
29.        map.type = MT_HIGH_VECTORS;
30.    #else
31.        map.type = MT_LOW_VECTORS;
```

```

32. #endif
33.     create_mapping(&map);
34.
35.     if (!vectors_high()) {
36.         map.virtual = 0;
37.         map.length = PAGE_SIZE * 2;
38.         map.type = MT_LOW_VECTORS;
39.         create_mapping(&map);
40.     }
41.
42.     /* Now create a kernel read-only mapping */
43.     map.pfn += 1;    devicemaps_init(mdesc);
44.
45.     map.virtual = 0xffff0000 + PAGE_SIZE;
46.     map.length = PAGE_SIZE;
47.     map.type = MT_LOW_VECTORS;
48.     create_mapping(&map);
49.
50.     .....
51. }

```

vectors = early_alloc(PAGE_SIZE * 2);

申请2 pages,从memory的bottom处分配。

early_trap_init(vectors);

```

1. void __init early_trap_init(void *vectors_base)
2. {
3.     #ifndef CONFIG_CPU_V7M
4.         unsigned long vectors = (unsigned long)vectors_base;
5.         extern char __stubs_start[], __stubs_end[];
6.         extern char __vectors_start[], __vectors_end[];
7.         unsigned i;
8.
9.         vectors_page = vectors_base;
10.
11.         /*
12.          * Poison the vectors page with an undefined instruction. This
13.          * instruction is chosen to be undefined for both ARM and Thumb
14.          * ISAs. The Thumb version is an undefined instruction with a
15.          * branch back to the undefined instruction.
16.          */
17.         for (i = 0; i < PAGE_SIZE / sizeof(u32); i++)
18.             ((u32 *)vectors_base)[i] = 0xe7fddef1;
19.
20.         /*
21.          * Copy the vectors, stubs and kuser helpers (in entry-armv.S)
22.          * into the vector page, mapped at 0xffff0000, and ensure these
23.          * are visible to the instruction stream.
24.          */
25.         memcpy((void *)vectors, __vectors_start, __vectors_end - __vectors_start)
26.         ;
27.         memcpy((void *)vectors + 0x1000, __stubs_start, __stubs_end - __stubs_start);
28.
29.         kuser_init(vectors_base);
30.
31.         flush_icache_range(vectors, vectors + PAGE_SIZE * 2);

```

```

31.         modify_domain(DOMAIN_USER, DOMAIN_CLIENT);
32.     #else /* ifndef CONFIG_CPU_V7M */
33.         /*
34.          * on V7-M there is no need to copy the vector table to a dedicated
35.          * memory area. The address is configurable and so a table in the kernel
36.          * image can be used.
37.          */
38.     #endif
39. }

```

```
extern char __stubs_start[], __stubs_end[];
```

```
extern char __vectors_start[], __vectors_end[];
```

定义在vmlinux.lds中

```

/*
 * The vectors and stubs are relocatable code, and the
 * only thing that matters is their relative offsets
 */

```

```
__vectors_start = .;
```

```
.vectors 0 : AT(__vectors_start) {
```

```
    *(.vectors)
```

```
}
```

```
. = __vectors_start + SIZEOF(.vectors);
```

```
__vectors_end = .;
```

```
__stubs_start = .;
```

```
.stubs 0x1000 : AT(__stubs_start) {
```

```

*(.stubs)

}

. = __stubs_start + SIZEOF(.stubs);

__stubs_end = .;

```

```
__vectors_start , __vectors_end
```

```
__stubs_start , __stubs_end
```

所代表的其实就是arch/arm/kernel/entry-armv.S中的最原始的trap handler。

整个entry-armv.S中的code分为两部分，分别放在2 sections，以便这两个section被mapping to 2 pages.

early_trap_init()的作用就是把entry-armv.S中的code copy到上面申请的2 pages中。

```

/*

* Create a mapping for the machine vectors at the high-vectors
* location (0xffff0000). If we aren't using high-vectors, also
* create a mapping at the low-vectors virtual address.

*/

map.pfn = __phys_to_pfn(virt_to_phys(vectors));

map.virtual = 0xffff0000;

map.length = PAGE_SIZE;

#ifdef CONFIG_KUSER_HELPERS

    map.type = MT_HIGH_VECTORS;

#else

```

```
map.type = MT_LOW_VECTORS;

#endif

create_mapping(&map);
```

接着就是把"vector" pages mapping to 0xffff0000。

```
if (!vectors_high()) {

    map.virtual = 0;

    map.length = PAGE_SIZE * 2;

    map.type = MT_LOW_VECTORS;

    create_mapping(&map);

}
```

如果ARM CPU并没有把vector page mapping to 0xffff,0000 , 则mapping to zero page。

```
#define vectors_high() (get_cr() & CR_V)

#define CR_V (1 << 13) /* Vectors relocated to 0xffff0000 */
```

```
/* Now create a kernel read-only mapping */

map.pfn += 1;

map.virtual = 0xffff0000 + PAGE_SIZE;

map.length = PAGE_SIZE;

map.type = MT_LOW_VECTORS;

create_mapping(&map);
```

接着mapping "stub" page。

