

- modify arch/arm/Makefile

```

1. walterzh@walterzh-Precision-T1650:~/work/current/tmp/work-shared/granite2/kernel-source/arch/arm$ diff Makefile ~/work2/original-kernel-src/linux-3.18.7/arch/arm/Makefile
2. 121,122c121,122
3. < KBUILD_CFLAGS +=$(CFLAGS_ABI) $(CFLAGS_ISA) $(arch-y) $(tune-y) $(call cc-option,-mshort-load-bytes,$(call cc-option,-malignment-traps,)) -mhard-float
   -Uarm
4. < KBUILD_AFLAGS +=$(CFLAGS_ABI) $(AFLAGS_ISA) $(arch-y) $(tune-y) -include asm/unified.h -mhard-float
5. ---
6. > KBUILD_CFLAGS +=$(CFLAGS_ABI) $(CFLAGS_ISA) $(arch-y) $(tune-y) $(call cc-option,-mshort-load-bytes,$(call cc-option,-malignment-traps,)) -msoft-float
   -Uarm
7. > KBUILD_AFLAGS +=$(CFLAGS_ABI) $(AFLAGS_ISA) $(arch-y) $(tune-y) -include asm/unified.h -msoft-float

```

convert `-msoft-float` option into `-mhard-float`

- enable kernel support VFP (NEON)
CONFIG_KERNEL_MODE_NEON=y



- How to do floating point computing in kernel code

```
1.  #include <linux/module.h>
2.  #include <linux/kernel.h>
3.  #include <linux/init.h>
4.  #include <asm/neon.h>      ①
5.
6.  static int __init test_module_init(void)
7.  {
8.      float d1, d2;
9.
10.     kernel_neon_begin();    ②
11.
12.     d1 = 1234.5;
13.     d2 = d1 / 6;
14.
15.     printk("1234.5 / 6 = %d\n", (int)d2);    ③
16.     kernel_neon_end();    ④
17.
18.     return 0;
19. }
20.
21. static void __exit test_module_exit(void)
22. {
23. }
24. module_init(test_module_init);
25. module_exit(test_module_exit);
26. MODULE_LICENSE("GPL");
```

```

1. walterzh@walterzh-Precision-T1650:~/work/victor/module-test/test$ cat Makefile
2. #
3. # =====
4. # Copyright (c) 2015 Marvell International, Ltd. All Rights Reserved
5. #
6. # Marvell Confidential
7. # =====
8. #
9.
10. obj-m := test.o
11.
12. ccflags-y += -O0 -march=armv7-a -mhard-float ⑤
13.
14. SRC := $(shell pwd)
15.
16. all:
17.     echo "test build"
18.     echo $(KERNEL_SRC)
19.     echo $(ccflags-y)
20.     $(MAKE) -C $(KERNEL_SRC) M=$(SRC)
21. modules_install:
22.     echo "install test"
23.     echo $(KERNEL_SRC)
24.     $(MAKE) -C $(KERNEL_SRC) M=$(SRC) modules_install
25.
26. clean:
27.     rm -f *.o *~ core *.depend *.cmd *.ko *.mod.c
28.     rm -f Module.markers Module.symvers modules.order
29.     rm -rf .tmp_versions Modules.symvers

```

①

You need include neon.h for accessing kernel_neon_begin() and kernel_neon_end() APIs

②④

All floating point computing must be embedded in kernel_neon_begin() and kernel_neon_end(), otherwise kernel will crash.

③

Because printk doesn't support "%f" specifier, so I convert float number into integer number to output.

⑤

You need add `-mhard-float` option to compile module

- internal of floating point computing

```

1. walterzh@walterzh-Precision-T1650:~/work/victor/module-test/test$ arm-linux-
gnueabi-objdump -dS test.ko
2.
3. test.ko:      file format elf32-littlearm
4.
5.
6. Disassembly of section .init.text:
7.
8. 00000000 <init_module>:
9. #include <linux/init.h>
10. #include <asm/neon.h>
11. #include <linux/preempt.h>
12.
13. static int __init test_module_init(void)
14. {
15.     0:   e52de004    push    {lr}          ; (str lr, [sp, #-4]!)
16.     4:   e24dd00c    sub    sp, sp, #12
17.         float d1, d2;
18.
19.         kernel_neon_begin();
20.     8:   ebfffffe    bl     0 <kernel_neon_begin>
21.
22.         d1 = 1234.5;
23.     c:   e3a03a05    mov    r3, #20480 ; 0x5000
24.    10:   e344349a    movt   r3, #17562 ; 0x449a
25.    14:   e58d3004    str    r3, [sp, #4]
26.         d2 = d1 / 6;
27.    18:   eddd7a01    vldr   s15, [sp, #4]          ⑥
28.    1c:   ed9f7a0c    vldr   s14, [pc, #48] ; 54 <init_module+0x54>
29.    20:   eec77a87    vdiv.f32 s15, s15, s14
30.    24:   edcd7a00    vstr   s15, [sp]
31.
32.         printk("1234.5 / 6 = %d\n", (int)d2);
33.    28:   eddd7a00    vldr   s15, [sp]          ⑦
34.    2c:   eefd7ae7    vcvtf.s32.f32 s15, s15
35.    30:   e3000000    movw   r0, #0
36.    34:   e3400000    movt   r0, #0
37.    38:   ee171a90    vmov   r1, s15
38.    3c:   ebfffffe    bl     0 <printk>
39.         kernel_neon_end();
40.    40:   ebfffffe    bl     0 <kernel_neon_end>
41.
42.         return 0;
43.    44:   e3a03000    mov    r3, #0
44.     }
45.    48:   e1a00003    mov    r0, r3
46.    4c:   e28dd00c    add    sp, sp, #12
47.    50:   e49df004    pop    {pc}          ; (ldr pc, [sp], #4)
48.    54:   40c00000    .word   0x40c00000
49.
50. Disassembly of section .exit.text:
51.
52. 00000000 <cleanup_module>:

```

```

53.
54. static void __exit test_module_exit(void)
55. {
56. }
57. 0: e12fff1e bx lr

```

⑥⑦

generate VFP(NEON) instruction

- Running result

```

1. root@granite2:~# insmod test.ko
2. 1234.5 / 6 = 205

```

- Restriction

floating point computing in kernel could run in the following context

1. non-interrupt context
2. preemption disable

That means you could not do floating point computing in interrupt service routine and you should not do time-consuming computing in kernel (Because preemption is disable, scheduler has been locked)

- The impact on kernel

Because `-mhard-float` option influenced deeply the generated kernel code, maybe we need more investigation to ensure whether the modification is safe.

for example:

in kernel/time/timeconv.c

```

1. /* do a mathdiv for long type */
2. static long math_div(long a, long b)
3. {
4.     return a / b - (a % b < 0);
5. }

```

if we use original gcc compiler option `-msoft-float`, the generated code as follow

```

1.  /* do a mathdiv for long type */
2.  static long math_div(long a, long b)
3.  {
4.      return a / b - (a % b < 0);
5.  c006b994:  e59d3010    ldr r3, [sp, #16]
6.  c006b998:  e0c10396    smull    r0, r1, r6, r3
7.  c006b99c:  e1a03fc6    asr r3, r6, #31
8.  c006b9a0:  e0862001    add r2, r6, r1
9.  c006b9a4:  e0633442    rsb r3, r3, r2, asr #8
10. c006b9a8:  e0832183    add r2, r3, r3, lsl #3
11. c006b9ac:  e0832182    add r2, r3, r2, lsl #3
12. c006b9b0:  e0822102    add r2, r2, r2, lsl #2
13. c006b9b4:  e0622006    rsb r2, r2, r6
14. c006b9b8:  e0433fa2    sub r3, r3, r2, lsr #31

```

Because the code has been optimized by `-O2`, it's not readable. But we could identify gcc doesn't generate VFP(NEON) instruction to implement "div" operation.

if we use gcc compiler option `-mhard-float`, the generated code as follow

```

1.  c0071fec <math_div>:
2.
3.  /* do a mathdiv for long type */
4.  static long math_div(long a, long b)
5.  {
6.  c0071fec:  e92d4070    push    {r4, r5, r6, lr}
7.  c0071ff0:  e1a06000    mov r6, r0
8.  c0071ff4:  e1a05001    mov r5, r1
9.      return a / b - (a % b < 0);
10. c0071ff8:  eb0ab4bb    bl  c031f2ec <__aeabi_idiv>
11. c0071ffc:  e1a04000    mov r4, r0
12. c0072000:  e1a00006    mov r0, r6
13. c0072004:  e1a01005    mov r1, r5
14. c0072008:  eb0ab51e    bl  c031f488 <__aeabi_idivmod>
15. }
16. c007200c:  e0440fa1    sub r0, r4, r1, lsr #31
17. c0072010:  e8bd8070    pop {r4, r5, r6, pc}

```

`__aeabi_idiv` and `__aeabi_idivmod` are introduced by "-mhard-float" option.

```
1. walterzh@walterzh-Precision-T1650:~/work/current/work/granite2-poky-linux-gn
ueabi/linux-granite-upstream/3.18.7+gitAUTOINC+26304af6aa-r0/image/boot$ arm
-linux-gnueabi-nm vmlinux-3.18.7-yocto-standard | grep "__aeabi_"
2. c031f2ec T __aeabi_idiv
3. c031f488 T __aeabi_idivmod
4. c031d52c T __aeabi_lasr
5. c031d510 T __aeabi_llsl
6. c031f4b0 T __aeabi_llsr
7. c031fce0 T __aeabi_lmul
8. c031f1b0 T __aeabi_udiv
9. c031f470 T __aeabi_udivmod
10. c031ff3c T __aeabi_ulcmp
11. c0015ce4 T __aeabi_unwind_cpp_pr0
12. c0015ce8 T __aeabi_unwind_cpp_pr1
13. c0015cec T __aeabi_unwind_cpp_pr2
```

These functions are from libgcc.a.

Amend:

We need not rebuild kernel by replacing `-mhard-float` option with `-msoft-float` .

The kernel is still VFP instruction free, built by `-msoft-float` option, but when build out-of-tree kernel module, we need modified arch/arm/Makefile.

So, do floating point computing in out-of-tree kernel module code is safe.