

in include/linux/init.h

/*

* A "pure" initcall has no dependencies on anything else, and purely

* initializes variables that couldn't be statically initialized.

*

* This only exists for built-in code, not for modules.

* Keep main.c:initcall_level_names[] in sync.

*/

#define pure_initcall(fn) __define_initcall(fn, 0)

#define core_initcall(fn) __define_initcall(fn, 1)

#define core_initcall_sync(fn) __define_initcall(fn, 1s)

#define postcore_initcall(fn) __define_initcall(fn, 2)

#define postcore_initcall_sync(fn) __define_initcall(fn, 2s)

#define arch_initcall(fn) __define_initcall(fn, 3)

#define arch_initcall_sync(fn) __define_initcall(fn, 3s)

#define subsys_initcall(fn) __define_initcall(fn, 4)

#define subsys_initcall_sync(fn) __define_initcall(fn, 4s)

#define fs_initcall(fn) __define_initcall(fn, 5)

#define fs_initcall_sync(fn) __define_initcall(fn, 5s)

#define rootfs_initcall(fn) __define_initcall(fn, rootfs)

#define device_initcall(fn) __define_initcall(fn, 6)

#define device_initcall_sync(fn) __define_initcall(fn, 6s)

```
#define late_initcall(fn)      __define_initcall(fn, 7)
```

```
#define late_initcall_sync(fn) __define_initcall(fn, 7s)
```

```
/* initcalls are now grouped by functionality into separate
```

```
 * subsections. Ordering inside the subsections is determined
```

```
 * by link order.
```

```
 * For backwards compatibility, initcall() puts the call in
```

```
 * the device init subsection.
```

```
 *
```

```
 * The `id' arg to __define_initcall() is needed so that multiple initcalls
```

```
 * can point at the same handler without causing duplicate-symbol build errors.
```

```
 */
```

```
#define __define_initcall(fn, id) \
```

```
    static initcall_t __initcall_###fn##id __used \
```

```
    __attribute__((__section__(".initcall" #id ".init"))) = fn; \
```

```
    LTO_REFERENCE_INITCALL(__initcall_###fn##id)
```

```
pure_initcall(ipc_ns_init);
```

```
==>
```

```
__define_initcall(ipc_ns_init, 0)
```

```
==>
```

```
#define __define_initcall(ipc_ns_init, 0)
```

```
static initcall_t __initcall_ipc_ns_init0 __used

__attribute__((__section__(".initcall0.init"))) = ipc_ns_init;
```

LTO_REFERENCE_INITCALL是为了fix 某些版本的gcc的bug，这里为NULL。

```
#ifdef CONFIG_LTO

/* Work around a LTO gcc problem: when there is no reference to a variable
 * in a module it will be moved to the end of the program. This causes
 * reordering of initcalls which the kernel does not like.
 * Add a dummy reference function to avoid this. The function is
 * deleted by the linker.
 */

#define LTO_REFERENCE_INITCALL(x) \
    ; /* yes this is needed */ \
    static __used __exit void *reference_##x(void) \
    { \
        return &x; \
    }

#else

#define LTO_REFERENCE_INITCALL(x)

#endif
```

定义了__initcall_ipc_ns_init0 variable,并放入section “**.initcall0.init**”.

而__define_initcall(XXX, 1) to __define_initcall(XXX, 7s)

定义了 __initcall_XXX_init1 , 放入section “**.initcall1.init**”.

__initcall_XXX_init1s,放入section “**.initcall1s.init**”.

__initcall_XXX_init2,放入section “**.initcall2.init**”.

__initcall_XXX_init2s,放入section “**.initcall2s.init**”.

.....

__initcall_XXX_init7,放入section “**.initcall7.init**”.

__initcall_XXX_init7s,放入section “**.initcall7s.init**”.

in vmlinux.lds中有

```
__initcall_start = .; *(.initcallearly.init) __initcall0_start = .; *(.initcall0.init) *(.initcall0s.init)
__initcall1_start = .; *(.initcall1.init) *(.initcall1s.init) __initcall2_start = .; *(.initcall2.init) *
(.initcall2s.init) __initcall3_start = .; *(.initcall3.init) *(.initcall3s.init) __initcall4_start = .; *(.initcall4.init)
*(.initcall4s.init) __initcall5_start = .; *(.initcall5.init) *(.initcall5s.init) __initcallrootfs_start = .; *
(.initcallrootfs.init) *(.initcallrootfss.init) __initcall6_start = .; *(.initcall6.init) *(.initcall6s.init)
__initcall7_start = .; *(.initcall7.init) *(.initcall7s.init) __initcall_end = .;
```

这些initcall function pointer组成了function pointer array。

在kernel initialization的后期会依次调用该array中的function pointer指向的function.

function被调用的先后次序完全由initcall level决定。

从最先被调用到最后调用依次为

level early [early_initcall(fn)]

level 0 [pure_initcall(fn)]

level 1 [core_initcall(fn)]

.....

in init/main.c

start_kernel()

|

|

\|

do_basic_setup()

|

|

\|

do_initcalls()

static void __init do_initcalls(void)

{

 int level;

 for (level = 0; level < ARRAY_SIZE(initcall_levels) - 1; level++)

 do_initcall_level(level);

}

static initcall_t *initcall_levels[] __initdata = {

 __initcall0_start,

```
__initcall1_start,  
__initcall2_start,  
__initcall3_start,  
__initcall4_start,  
__initcall5_start,  
__initcall6_start,  
__initcall7_start,  
__initcall_end,  
};
```

```
static void __init do_initcall_level(int level)  
{  
    initcall_t *fn;  
  
    strcpy(initcall_command_line, saved_command_line);  
  
    parse_args(initcall_level_names[level],  
               initcall_command_line, __start__param,  
               __stop__param - __start__param,  
               level, level,  
               &repair_env_string);  
  
    for (fn = initcall_levels[level]; fn < initcall_levels[level+1]; fn++)  
        do_one_initcall(*fn);  
}
```

```
int __init_or_module do_one_initcall(initcall_t fn)
{
    int count = preempt_count();

    int ret;

    char msgbuf[64];

    if (initcall_blacklisted(fn))          (1)
        return -EPERM;

    if (initcall_debug)                    (2)
        ret = do_one_initcall_debug(fn);
    else
        ret = fn();

    msgbuf[0] = 0;

    if (preempt_count() != count) {
        sprintf(msgbuf, "preemption imbalance ");
        preempt_count_set(count);
    }

    if (irqs_disabled()) {
        strcat(msgbuf, "disabled interrupts ", sizeof(msgbuf));

        local_irq_enable();
    }

    WARN(msgbuf[0], "initcall %pF returned with %s\n", fn, msgbuf);
}
```

```
return ret;  
  
}
```

(1)

`initcall_blacklist=` [KNL] Do not execute a comma-separated list of
initcall functions. Useful for debugging built-in
modules and initcalls.

(2)

`initcall_debug` [KNL] Trace initcalls as they are executed. Useful
for working out where the kernel is dying during
startup.

built-in driver的init function都是在do_initcalls()阶段执行的,也就是各个driver的probe() function在此时执行。

通过“initcall_debug” kernel parameter可以dump出driver初始化的顺序,对调试built-in driver的initialization很有利。

而“initcall_blacklist” kernel parameter则可以暂时禁止某个driver。因为不执行该driver的init function,就不会调用该driver的probe(),自然该driver也不工作了。

另外如果各个driver之间有依赖关系,那么通过安排不同的initcall level,可以建立正确的dependency.

for example,

GPIO即是普通device，同时也可作为interrupt controller,而某些device的interrupt line挂在GPIO的interrupt controller上，那么developer必须保证GPIO

先于该设备初始化，否则该设备初始化必然失败。因为一般request IRQ都在driver的probe()中，而如果GPIO driver还未初始化好，那么该device的driver request IRQ必然失败。