

Virtual kernel memory layout:

vector : 0xffff0000 - 0xffff1000 ( 4 kB)

fixmap : 0xffc00000 - 0xffe00000 (2048 kB)

vmalloc : 0xf0000000 - 0xff000000 ( 240 MB)

lowmem : 0xc0000000 - 0xef800000 ( 760 MB)

pkmap : 0xbfe00000 - 0xc0000000 ( 2 MB)

**modules : 0xbf000000 - 0xbfe00000 ( 14 MB)**

.text : 0xc0008000 - 0xc05cde80 (5912 kB)

.init : 0xc05ce000 - 0xc0602000 ( 208 kB)

.data : 0xc0602000 - 0xc0638728 ( 218 kB)

.bss : 0xc0638728 - 0xc06a8af4 ( 449 kB)

kernel module被载入(load)到[0xbf000000, 0xbfe00000)的14M空间。

```
root@granite2:~# cat /proc/modules
```

```
galcore 160452 0 - Live 0xbf245000 (O)
```

```
ipv6 276100 20 [permanent], Live 0xbf1eb000
```

```
imagepower 1782 0 - Live 0xbf1e7000 (O)
```

```
stepper_api_a0 12747 0 - Live 0xbf1df000 (O)
```

```
upc 88464 0 - Live 0xbf1ab000 (O)
```

```
pegmatite_regulator 4051 4 imagepower,upc,[permanent], Live 0xbf1a3000
```

```
picdriver 55024 0 - Live 0xbf175000 (O)
```

```
cisxdriver 34852 0 - Live 0xbf153000 (O)
```

```
laservideo_a0 80001 0 - Live 0xbf135000 (O)
```

```
scanblkdriver 55179 0 - Live 0xbf11c000 (O)
```

```
icetestdriver 15129 0 - Live 0xbf113000 (O)
```

```
pieedriver 179267 0 - Live 0xbf0d7000 (O)
```

```
dmaalloc 4127 1 laservideo_a0, Live 0xbf0cc000 (O)
```

```
dros 9786 2 laservideo_a0,scanblkdriver, Live 0xbf0c4000 (O)
```

```
hips_pll 4966 1 laservideo_a0, Live 0xbf0ac000 (O)
```

```
mv61_cdma 31142 0 - Live 0xbf09f000
```

```
sccplite 5695 0 - Live 0xbf093000 (O)
```

```
stepper_mod_a0 15653 1 stepper_api_a0, Live 0xbf05b000 (O)
```

```
ehci_hcd 44353 0 - Live 0xbf04b000
```

```
ci_hdrc_imx 2976 0 - Live 0xbf047000
```

```
usbmisc_imx 4330 1 ci_hdrc_imx, Live 0xbf042000
```

```
ci_hdrc 20768 1 ci_hdrc_imx, Live 0xbf038000
```

```
ipc_driver 4578 0 - Live 0xbf010000 (O)
```

```
i2c_pxa 8241 0 - Live 0xbf009000
```

dro\_pegmatite 1598 0 - Live 0xbf005000

pegmatite\_wdt 4332 0 - Live 0xbf000000

**galcore 160452 0 - Live 0xbf245000 (O)**

galcore载入的地址为0xbf245000,size为160452 bytes.

当kernel载入一个module时, 分配memory

kernel/module.c

load\_module() --> move\_module() --> module\_alloc\_update\_bounds() --> module\_alloc(size)

```
1.  static void *module_alloc_update_bounds(unsigned long size)
2.  {
3.      void *ret = module_alloc(size);
4.
5.      if (ret) {
6.          mutex_lock(&module_mutex);
7.          /* Update module bounds. */
8.          if ((unsigned long)ret < module_addr_min)
9.              module_addr_min = (unsigned long)ret;
10.         if ((unsigned long)ret + size > module_addr_max)
11.             module_addr_max = (unsigned long)ret + size;
12.         mutex_unlock(&module_mutex);
13.     }
14.     return ret;
15. }
```

所有module占用的空间在[module\_addr\_min, module\_addr\_max)之间。

/\* Bounds of module allocation, for speeding \_\_module\_address.

\* Protected by module\_mutex. \*/

static unsigned long module\_addr\_min = -1UL, module\_addr\_max = 0;

查看以下这两个variables的值。

```
walterzh$ nm vmlinux-3.18.7-yocto-standard | grep module_addr_min
```

```
c061268c d module_addr_min
```

```
walterzh$ nm vmlinux-3.18.7-yocto-standard | grep module_addr_max
```

```
c068d288 b module_addr_max
```

```
root@granite2:~# devmem 0x0061268c
```

```
0xBF000000
```

```
root@granite2:~# devmem 0x0061268c 8d288
```

```
0xBF276C24
```

[0xBF000000 , 0xBF276C24 ) 之间。

```
1. void * __weak module_alloc(unsigned long size)
2. {
3.     return vmalloc_exec(size);
4. }
```

```

1.  /**
2.   *      vmalloc_exec - allocate virtually contiguous, executable memory
3.   *      @size:          allocation size
4.   *
5.   *      Kernel-internal function to allocate enough pages to cover @size
6.   *      the page level allocator and map them into contiguous and
7.   *      executable kernel virtual space.
8.   *
9.   *      For tight control over page level allocator and protection flags
10.   *      use __vmalloc() instead.
11.   */
12.
13. void *vmalloc_exec(unsigned long size)
14. {
15.     return __vmalloc_node(size, 1, GFP_KERNEL | __GFP_HIGHMEM, PAGE_KERNEL_EXEC,
16.                           C,
17.                           NUMA_NO_NODE, __builtin_return_address(0));

```

module占用的空间优先从highmem分配。

并且由于module是code，所以占用的空间必须是executable。

```

1.  static void *__vmalloc_node(unsigned long size, unsigned long align,
2.                             gfp_t gfp_mask, pgprot_t prot,
3.                             int node, const void *caller)
4.  {
5.      return __vmalloc_node_range(size, align, VMALLOC_START, VMALLOC_END,
6.                                  gfp_mask, prot, node, caller);
7.  }

```

```
#ifndef CONFIG_THUMB2_KERNEL
```

```
#define MODULES_VADDR      (PAGE_OFFSET - SZ_16M)
```

```
#else
```

```
/* smaller range for Thumb-2 symbols relocation (2^24)*/
```

```
#define MODULES_VADDR      (PAGE_OFFSET - SZ_8M)
```

```
#endif
```

```
/*
```

```
* The highmem pkmap virtual space shares the end of the module area.
```

```
*/
```

```
#ifdef CONFIG_HIGHMEM
```

```
#define MODULES_END          (PAGE_OFFSET - PMD_SIZE)
```

```
#else
```

```
#define MODULES_END          (PAGE_OFFSET)
```

```
#endif
```

```
#define PMD_SHIFT            21
```

```
#define PMD_SIZE             (1UL << PMD_SHIFT)
```

```
MODULES_VADDR = 0xc0000000 - 16M = 0xbf000000
```

```
MODULES_END = 0xc0000000 - 1UL << 21 = 0xc0000000 - 2M = 0xbfe00000
```

We could get the verification from the following boot log.

Virtual kernel memory layout:

vector : 0xffff0000 - 0xffff1000 ( 4 kB)

fixmap : 0xffc00000 - 0xffe00000 (2048 kB)

vmalloc : 0xf0000000 - 0xff000000 ( 240 MB)

lowmem : 0xc0000000 - 0xef800000 ( 760 MB)

pkmap : 0xbfe00000 - 0xc0000000 ( 2 MB)

**modules : 0xbf000000 - 0xbfe00000 ( 14 MB)**

.text : 0xc0008000 - 0xc05cde60 (5912 kB)

.init : 0xc05ce000 - 0xc0602000 ( 208 kB)

.data : 0xc0602000 - 0xc0638728 ( 218 kB)

.bss : 0xc0638728 - 0xc06a8af4 ( 449 kB)

/\*\*

- \* \_\_vmalloc\_node\_range - allocate virtually contiguous memory
- \* @size: allocation size
- \* @align: desired alignment
- \* @start: vm area range start
- \* @end: vm area range end
- \* @gfp\_mask: flags for the page level allocator
- \* @prot: protection mask for the allocated pages
- \* @node: node to use for allocation or NUMA\_NO\_NODE
- \* @caller: caller's return address
- \*
- \* Allocate enough pages to cover @size from the page level
- \* allocator with @gfp\_mask flags. Map them into contiguous
- \* kernel virtual space, using a pagetable protection of @prot.

\*/

```
void *__vmalloc_node_range(unsigned long size, unsigned long align,  
                           unsigned long start, unsigned long end, gfp_t gfp_mask,  
                           pgprot_t prot, int node, const void *caller);
```

```
static void *__vmalloc_node(unsigned long size, unsigned long align,  
                             gfp_t gfp_mask, pgprot_t prot,  
                             int node, const void *caller)  
{  
    return __vmalloc_node_range(size, 1, 0xbf000000, 0xbfe00000,  
                                GFP_KERNEL | __GFP_HIGHMEM,  
                                PAGE_KERNEL_EXEC,  
                                NUMA_NO_NODE,  
                                __builtin_return_address(0));
```

最终在alloc\_vmap\_area() / vmalloc.c中在[0xbf000000, 0xbfe00000)的virtual address space中分配一块size大小的

空间 ( virtual address is continuous, but physical address maybe are discrete. )



