

start\_kernel()的最后是rest\_init()

```
1.  asmlinkage __visible void __init start_kernel(void)
2.  {
3.
4.  .....
5.
6.  /* Do the rest non-__init'ed, we're now alive */
7.      rest_init();
8.  }
```

```

1. static __initdata DECLARE_COMPLETION(kthreadd_done);
2.
3. static noinline void __init_refok rest_init(void)
4. {
5.     int pid;
6.
7.     rcu_scheduler_starting();
8.     /*
9.      * We need to spawn init first so that it obtains pid 1, however
10.     * the init task will end up wanting to create kthreads, which, if
11.     * we schedule it before we create kthreadd, will OOPS.
12.     */
13.     kernel_thread(kernel_init, NULL, CLONE_FS);      (0)
14.     numa_default_policy();                          (0.01)
15.     pid = kernel_thread(kthreadd, NULL, CLONE_FS | CLONE_FILES);
16.     rcu_read_lock();
17.     kthreadd_task = find_task_by_pid_ns(pid, &init_pid_ns);
18.     rcu_read_unlock();
19.     complete(&kthreadd_done);                      (1)
20.
21.     /*
22.     * The boot idle thread must execute schedule()
23.     * at least once to get things moving:
24.     */
25.     init_idle_bootup_task(current);
26.     schedule_preempt_start_kernel();的最后是rest_init()

```

```

asmlinkage __visible void __init start_kernel(void)
{
    .....

    /* Do the rest non-__init'ed, we're now alive */
    rest_init();
}_disabled();

```

```
27.  /* Call into cpu_idle with preempt disabled */
28.      cpu_startup_entry(CPUHP_ONLINE);
29.  }
```

static int \_\_ref kernel\_init(void \*unused)和

int kthreadd(void \*unused)

分别在独立的thread中运行，但两者间还是有同步关系的。

```

1. static int __ref kernel_init(void *unused)
2. {
3.     int ret;
4.
5.     kernel_init_freeable();           (0.1)
6.     /* need to finish all async __init code before freeing the memory */
7.     async_synchronize_full();
8.     free_initmem();
9.     mark_rodata_ro();
10.    system_state = SYSTEM_RUNNING;
11.    numa_default_policy();
12.
13.    flush_delayed_fput();
14.
15.    if (ramdisk_execute_command) {
16.        ret = run_init_process(ramdisk_execute_command);
17.        if (!ret)
18.            return 0;
19.        pr_err("Failed to execute %s (error %d)\n",
20.            ramdisk_execute_command, ret);
21.    }
22.
23.    /*
24.     * We try each of these until one succeeds.
25.     *
26.     * The Bourne shell can be used instead of init if we are
27.     * trying to recover a really broken machine.
28.     */
29.    if (execute_command) {             (2)
30.        ret = run_init_process(execute_command);
31.        if (!ret)

```

```
32.         return 0;
33.         pr_err("Failed to execute %s (error %d).  Attempting defaults...\n",
34.               execute_command, ret);
35.     }
36.     if (!try_to_run_init_process("/sbin/init") ||
37.         !try_to_run_init_process("/etc/init") ||
38.         !try_to_run_init_process("/bin/init") ||
39.         !try_to_run_init_process("/bin/sh"))
40.         return 0;
41.
42.     panic("No working init found.  Try passing init= option to kernel.  "
43.          "See Linux Documentation/init.txt for guidance.");
44. }
```

```
1. static noinline void __init kernel_init_freeable(void)
2. {
3.     /*
4.      * Wait until kthreadd is all set-up.
5.      */
6.     wait_for_completion(&kthreadd_done);           (0.2)
7.
8.     /* Now the scheduler is fully set up and can do blocking allocations */
9.     gfp_allowed_mask = __GFP_BITS_MASK;           (1.1)
10.
11.    /*
12.     * init can allocate pages on any node
13.     */
14.    set_mems_allowed(node_states[N_MEMORY]);
15.    /*
16.     * init can run on any cpu.
17.     */
18.    set_cpus_allowed_ptr(current, cpu_all_mask);
19.
20.    cad_pid = task_pid(current);
21.
22.    smp_prepare_cpus(setup_max_cpus);
23.
24.    do_pre_smp_initcalls();
25.    lockup_detector_init();
26.
27.    smp_init();
28.    sched_init_smp();
29.
30.    do_basic_setup();                               (1.2)
31.
```

```

32.      /* Open the /dev/console on the rootfs, this should never fail */
33.      if (sys_open((const char __user *) "/dev/console", O_RDWR, 0) < 0)
34.          pr_err("Warning: unable to open an initial console.\n");
35.
36.      (void) sys_dup(0);
37.      (void) sys_dup(0);
38.      /*
39.       * check if there is an early userspace init.  If yes, let it do all
40.       * the work
41.       */
42.
43.      if (!ramdisk_execute_command)
44.          ramdisk_execute_command = "/init";
45.
46.      if (sys_access((const char __user *) ramdisk_execute_command, 0) != 0) {
47.          ramdisk_execute_command = NULL;
48.          prepare_namespace();
49.      }
50.
51.      /*
52.       * Ok, we have completed the initial bootup, and
53.       * we're essentially up and running. Get rid of the
54.       * initmem segments and start the user-mode stuff..
55.       */
56.
57.      /* rootfs is available now, try loading default modules */
58.      load_default_modules();
59.  }

```

( 0 ) \_\_\_\_\_ (0.1) \_\_\_\_\_ (0.2) \_\_\_\_\_ (1) \_\_\_\_\_ (1.1) \_\_\_\_\_ (1.2) \_\_\_\_\_ (2)

| \_\_\_\_\_ (0.01) \_\_\_\_\_ (1)

所以create init process (2)总在driver初始化(1.2)之后。