

Memory: 1028396K/1044480K available (4482K kernel code, 190K rwd data, 1280K rodata, 204K init, 448K bss, 16084K reserved, 270336K highmem)

Virtual kernel memory layout:

vector : 0xffff0000 - 0xffff1000 ( 4 kB)

fixmap : 0xffc00000 - 0xffe00000 (2048 kB)

vmalloc : 0xf0000000 - 0xff000000 ( 240 MB)

lowmem : 0xc0000000 - 0xef800000 ( 760 MB)

pkmap : 0xbfe00000 - 0xc0000000 ( 2 MB)

modules : 0xbf000000 - 0xbfe00000 ( 14 MB)

.text : 0xc0008000 - 0xc05a8d84 (5764 kB)

.init : 0xc05a9000 - 0xc05dc000 ( 204 kB)

.data : 0xc05dc000 - 0xc060b980 ( 191 kB)

.bss : 0xc060b980 - 0xc067bd34 ( 449 kB)

## 0. 对low mamory而言 ( )

`paddr = vaddr - PAGE_OFFSET`

## 1. vmalloc 非连续空间的物理映射, VMALLOC\_START to VMALLOC\_END

## 2. pkmap空间

```
#define PKMAP_BASE (PAGE_OFFSET - PMD_SIZE)
```

2M size

在x86体系结构上，高于896MB的所有物理内存的范围大都是高端内存

arch/arm/highmem.c

\*

\* Author: Nicolas Pitre

\* Created: september 8, 2008

\* Copyright: Marvell Semiconductors Inc.

\*

kmap/unkmap系统调用是用来映射高端物理内存页到内核地址空间的api函数，他们分配的内存虚拟地址范围属于[PKMAP\_BASE, PAGE\_OFFSET]即[0xbfe00000, 0xc0000000]范围，大小是2M的虚拟空间，为了映射该块虚拟地址，所使用的二级页表的大小刚好是一个物理page的总计是两个pte table ( 4KB )

```
void *kmap(struct page *page)
```

```

{
    might_sleep();

    if (!PageHighMem(page))
        return page_address(page);

    return kmap_high(page);
}

EXPORT_SYMBOL(kmap);

```

if physical page is highmem, virtual address = kmap\_high(page) in pkmap region.

```

271 /**
272  * kmap_high - map a highmem page into memory
273  * @page: &struct page to map
274  *
275  * Returns the page's virtual memory address.
276  *
277  * We cannot call this from interrupts, as it may block.
278  */
279 void *kmap_high(struct page *page)
280 {
281     unsigned long vaddr;
282
283     /*
284      * For highmem pages, we can't trust "virtual" until
285      * after we have the lock.

```

```
286     */
287     lock_kmap();
288     vaddr = (unsigned long)page_address(page);
289     if (!vaddr)    为空，表示还未建立virtual-physical mapping
290         vaddr = map_new_virtual(page);
291     pkmap_count[PKMAP_NR(vaddr)]++;
292     BUG_ON(pkmap_count[PKMAP_NR(vaddr)] < 2);
293     unlock_kmap();
294     return (void*) vaddr;
295 }
```

```
217 static inline unsigned long map_new_virtual(struct page *page)
```

```
218 {
219     unsigned long vaddr;
220     int count;
221     unsigned int last_pkmap_nr;
222     unsigned int color = get_pkmap_color(page);
223
224 start:
225     count = get_pkmap_entries_count(color);
226     /* Find an empty entry */
227     for (;;) {
228         last_pkmap_nr = get_next_pkmap_nr(color);
229         if (no_more_pkmaps(last_pkmap_nr, color)) {
230             flush_all_zero_pkmaps();
```

```
231         count = get_pkmap_entries_count(color);
232     }
233     if (!pkmap_count[last_pkmap_nr])
234         break; /* Found a usable entry */
235     if (--count)
236         continue;
237
238     /*
239     * Sleep for somebody else to unmap their entries
240     */
241     {
242         DECLARE_WAITQUEUE(wait, current);
243         wait_queue_head_t *pkmap_map_wait =
244             get_pkmap_wait_queue_head(color);
245
246         __set_current_state(TASK_UNINTERRUPTIBLE);
247         add_wait_queue(pkmap_map_wait, &wait);
248         unlock_kmap();
249         schedule();
250         remove_wait_queue(pkmap_map_wait, &wait);
251         lock_kmap();
252
253         /* Somebody else might have mapped it while we slept */
254         if (page_address(page))
255             return (unsigned long)page_address(page);
```

```

256
257         /* Re-start */
258         goto start;
259     }
260 }
261 vaddr = PKMAP_ADDR(last_pmap_nr);
262 set_pte_at(&init_mm, vaddr,      建立virtual-physical之间的mapping
263          &(pmap_page_table[last_pmap_nr]), mk_pte(page, kmap_prot));
264
265 pmap_count[last_pmap_nr] = 1;
266 set_page_address(page, (void *)vaddr);
267
268 return vaddr;
269 }

```

```

#define PKMAP_ADDR(nr)      (PKMAP_BASE + ((nr) << PAGE_SHIFT))

```

### 3. module载入空间

modules : 0xbf000000 - 0xbfe00000 ( 14 MB)

```
root@granite2:~# cat /proc/modules
```

```
galcore 161066 0 - Live 0xbf27f000 (O)
```

```
ipv6 276100 20 [permanent], Live 0xbf225000
```

```
imagepower 1865 0 - Live 0xbf221000 (O)
```

```
laservideo_a0 80001 0 - Live 0xbf203000 (O)
```

upc 88464 0 - Live 0xbf1c9000 (O)

icetestdriver 15129 0 - Live 0xbf1c0000 (O)

scanblkdriver 52888 0 - Live 0xbf1a8000 (O)

picdriver 55319 0 - Live 0xbf17a000 (O)

dmaalloc 4239 1 laservideo\_a0, Live 0xbf175000 (O)

dros 9633 1 laservideo\_a0, Live 0xbf16d000 (O)

stepper\_api\_b0 12711 0 - Live 0xbf165000 (O)

pieedriver 179267 0 - Live 0xbf129000 (O)

cisxdriver 34852 0 - Live 0xbf111000 (O)

pegmatite\_regulator 4390 4 imagepower,upc,[permanent], Live 0xbf100000

hips\_pll 4966 1 laservideo\_a0, Live 0xbf0cd000 (O)

m25p80 6875 0 - Live 0xbf0c7000

mv61\_cdma 31142 0 - Live 0xbf0ba000

spi\_nor 13262 1 m25p80, Live 0xbf0b2000

stepper\_mod\_b0 16545 1 stepper\_api\_b0, Live 0xbf076000 (O)

sccplite 5695 0 - Live 0xbf06a000 (O)

ehci\_hcd 44353 0 - Live 0xbf048000

spi\_pxa2xx\_platform 14149 0 - Live 0xbf040000

dcmotor\_reg 7724 0 - Live 0xbf028000 (O)

ipc\_driver 4578 0 - Live 0xbf023000 (O)

i2c\_pxa 7548 0 - Live 0xbf005000

pegmatite\_wdt 4332 0 - Live 0xbf000000

#### 4. 总体memory map

.....

Built 1 zonelists in Zone order, mobility grouping on. **Total pages: 259600**

Kernel command line: console=ttyS0,115200n8 earlyprintk=serial,ttyS0,115200  
root=/dev/mmcblk1p2 uio\_pdrv\_genirq.of\_id=generic-uio rootwait

PID hash table entries: 4096 (order: 2, 16384 bytes)

Dentry cache hash table entries: 131072 (order: 7, 524288 bytes)

Inode-cache hash table entries: 65536 (order: 6, 262144 bytes)

Memory: 1028396K/1044480K available (4482K kernel code, 190K rwdma, 1280K rodata, 204K init,  
448K bss, 16084K reserved, **270336K highmem**)

lowmem : 0xc0000000 - 0xef800000 ( 760 MB)

270336K highmem

760 MB + 270M = 1030M

pkmap : 0xbfe00000 - 0xc0000000 ( 2 MB)

要管理270M high mem ?

内核一次最多只能管理2MB的high mem。超过之，则需要临时unmap原来的，mapping新的。



