

in mm/memblock.c

Use global variable memblock to track all available **physical** memory and all reserved **physical** memory.

```
struct memblock memblock __initdata_memblock = {  
  
    .memory.regions      = memblock_memory_init_regions,  
  
    .memory.cnt          = 1, /* empty dummy entry */  
  
    .memory.max          = INIT_MEMBLOCK_REGIONS,  
  
  
    .reserved.regions    = memblock_reserved_init_regions,  
  
    .reserved.cnt        = 1, /* empty dummy entry */  
  
    .reserved.max        = INIT_MEMBLOCK_REGIONS,  
  
  
  
  
#ifdef CONFIG_HAVE_MEMBLOCK_PHYS_MAP  
  
    .physmem.regions     = memblock_physmem_init_regions,  
  
    .physmem.cnt         = 1, /* empty dummy entry */  
  
    .physmem.max         = INIT_PHYSMEM_REGIONS,  
  
#endif  
  
  
  
    .bottom_up           = false,  
  
    .current_limit       = MEMBLOCK_ALLOC_ANYWHERE,  
  
};
```

CONFIG_HAVE_MEMBLOCK_PHYS_MAP is not set in Granite2 and Gemstone2 LSP.

```
struct memblock {  
  
    bool bottom_up; /* is bottom up direction? */  
  
    phys_addr_t current_limit;
```

struct memblock_type memory; ①

struct memblock_type reserved; ②

#ifdef CONFIG_HAVE_MEMBLOCK_PHYS_MAP

struct memblock_type physmem;

#endif

};

①

record all available physical memory in it (array)

②

record all reserved physical memory in it (array)

struct memblock_type {

unsigned long cnt; /* number of regions */

unsigned long max; /* size of the allocated array */

phys_addr_t total_size; /* size of all regions */

struct memblock_region *regions;

};

struct memblock_region {

phys_addr_t base;

phys_addr_t size;

unsigned long flags;

#ifdef CONFIG_HAVE_MEMBLOCK_NODE_MAP

int nid;

#endif

};

```
static struct memblock_region memblock_memory_init_regions[INIT_MEMBLOCK_REGIONS]
__initdata_memblock;

static struct memblock_region memblock_reserved_init_regions[INIT_MEMBLOCK_REGIONS]
__initdata_memblock;
```

```
#define INIT_MEMBLOCK_REGIONS 128
```

即无论是"memory"的memblock_type还是"reserved"的memblock_type,最多可以有128个region。

```
global variable memblock.memory.regions = memblock_memory_init_regions;

global variable memblock.reserved.regions = memblock_reserved_init_regions;
```

```
=====
```

in include/linux/memblock.h

```
int memblock_add(phys_addr_t base, phys_addr_t size);           ①

int memblock_reserve(phys_addr_t base, phys_addr_t size);      ②
```

①

把[base, base + size)的region添加到"memory"的memblock_type(memblock.memory.regions)中去

"memory"的memblock_type中的region是Linux可以使用的total memory(包括了下面"reserved"的memblock_type中的memory)

②

把[base, base + size)的region添加到"reserved"的memblock_type(memblock.reserved.regions)中去

"reserved"的memblock_type中的region则是已经被占用，不能被分配出去的memory

```
=====
```

About memblock_add()

memblock_add()用于告知Linux可用的memory。目前有两种方式向Linux kernel报告可用memory.

1. atag方式，即u-boot通过ATAG_MEM这个tag来告知kernel当前系统有那些memory(通过ATAG_MEM)。这是比较老的方式，使用device tree以后就不适用了。

2. 在dts中指定memory。这是device tree的方式。

比如在mv6270-toc.dts中

```
memory {  
    device_type = "memory";  
    reg = <0 0x00000000 0 0x40000000>; /* 1 GB */  
};
```

该device node就用于指定系统的memory。

in arch/arm/kernel/setup.c

```
void __init setup_arch(char **cmdline_p)  
{  
    const struct machine_desc *mdesc;  
  
    setup_processor();  
  
    mdesc = setup_machine_fdt(__atags_pointer);  
  
    if (!mdesc)          (A)
```

```

        mdesc = setup_machine_tags(__atags_pointer, __machine_arch_type);

machine_desc = mdesc;

machine_name = mdesc->name;

.....

}

```

(A)

如果mdesc为NULL，表示kernel没找到device tree对应的dtb，则寻找ATAG(老的方式)。

in arch/arm/kernel/devtree.c

```

/**
 * setup_machine_fdt - Machine setup when an dtb was passed to the kernel
 * @dt_phys: physical address of dt blob
 *
 * If a dtb was passed to the kernel in r2, then use it to choose the
 * correct machine_desc and to setup the system.
 */
const struct machine_desc * __init setup_machine_fdt(unsigned int dt_phys)
{
    const struct machine_desc *mdesc, *mdesc_best = NULL;

#ifdef CONFIG_ARCH_MULTIPLATFORM
    DT_MACHINE_START(GENERIC_DT, "Generic DT based system")

    MACHINE_END

    mdesc_best = &__mach_desc_GENERIC_DT;
#endif
}

```

```

if (!dt_phys || !early_init_dt_verify(phys_to_virt(dt_phys)))

    return NULL;

mdesc = of_flat_dt_match_machine(mdesc_best, arch_get_next_mach);

if (!mdesc) {

    const char *prop;

    int size;

    unsigned long dt_root;

    early_print("\nError: unrecognized/unsupported "

        "device tree compatible list:\n[ ");

    dt_root = of_get_flat_dt_root();

    prop = of_get_flat_dt_prop(dt_root, "compatible", &size);

    while (size > 0) {

        early_print("%s' ", prop);

        size -= strlen(prop) + 1;

        prop += strlen(prop) + 1;

    }

    early_print("]\n\n");

    dump_machine_table(); /* does not return */

}

/* We really don't want to do this, but sometimes firmware provides buggy data */

if (mdesc->dt_fixup)

    mdesc->dt_fixup();

```

```

early_init_dt_scan_nodes();

/* Change machine number to match the mdesc we're using */
__machine_arch_type = mdesc->nr;

return mdesc;
}

void __init early_init_dt_scan_nodes(void)
{
    /* Retrieve various information from the /chosen node */
    of_scan_flat_dt(early_init_dt_scan_chosen, boot_command_line);

    /* Initialize {size,address}-cells info */
    of_scan_flat_dt(early_init_dt_scan_root, NULL);

    /* Setup memory, calling early_init_dt_add_memory_arch */
    of_scan_flat_dt(early_init_dt_scan_memory, NULL);
}

/**
 * early_init_dt_scan_memory - Look for an parse memory nodes
 */
int __init early_init_dt_scan_memory(unsigned long node, const char *uname,
                                     int depth, void *data)
{
    const char *type = of_get_flat_dt_prop(node, "device_type", NULL);

    const __be32 *reg, *endp;

```

```
int l;
```

```
/* We are scanning "memory" nodes only */
```

```
if (type == NULL) {
```

```
    /*
```

```
    * The longtrail doesn't have a device_type on the
```

```
    * /memory node, so look for the node called /memory@0.
```

```
    */
```

```
    if (!IS_ENABLED(CONFIG_PPC32) || depth != 1 || strcmp(uname, "memory@0") != 0)
```

```
        return 0;
```

```
} else if (strcmp(type, "memory") != 0)
```

④

```
    return 0;
```

```
reg = of_get_flat_dt_prop(node, "linux,usable-memory", &l);
```

```
if (reg == NULL)
```

```
    reg = of_get_flat_dt_prop(node, "reg", &l);
```

⑤

```
if (reg == NULL)
```

```
    return 0;
```

```
endp = reg + (l / sizeof(__be32));
```

```
pr_debug("memory scan node %s, reg size %d, data: %x %x %x %x\n",
```

```
    uname, l, reg[0], reg[1], reg[2], reg[3]);
```

```
while ((endp - reg) >= (dt_root_addr_cells + dt_root_size_cells)) {
```

```
    u64 base, size;
```

```
    base = dt_mem_next_cell(dt_root_addr_cells, &reg);
```

```
    size = dt_mem_next_cell(dt_root_size_cells, &reg);
```



```

if (size == 0)

    continue;

pr_debug(" - %llx , %llx\n", (unsigned long long)base,

    (unsigned long long)size);

```

```

early_init_dt_add_memory_arch(base, size);

```

⑥

```

}

```

```

return 0;

```

```

}

```

in drivers/of/fdt.c

```

void __init __weak early_init_dt_add_memory_arch(u64 base, u64 size)

```

```

{

```

```

    const u64 phys_offset = __pa(PAGE_OFFSET);

```

```

    if (!PAGE_ALIGNED(base)) {

```

```

        size -= PAGE_SIZE - (base & ~PAGE_MASK);

```

```

        base = PAGE_ALIGN(base);

```

```

    }

```

```

    size &= PAGE_MASK;

```

```

    if (base > MAX_PHYS_ADDR) {

```

```

        pr_warning("Ignoring memory block 0x%llx - 0x%llx\n",

```

```

            base, base + size);

```

```

        return;

```

```

    }

```

```

if (base + size - 1 > MAX_PHYS_ADDR) {

    pr_warning("Ignoring memory range 0x%llx - 0x%llx\n",
               ((u64)MAX_PHYS_ADDR) + 1, base + size);

    size = MAX_PHYS_ADDR - base + 1;

}

if (base + size < phys_offset) {

    pr_warning("Ignoring memory block 0x%llx - 0x%llx\n",
               base, base + size);

    return;

}

if (base < phys_offset) {

    pr_warning("Ignoring memory range 0x%llx - 0x%llx\n",
               base, phys_offset);

    size -= phys_offset - base;

    base = phys_offset;

}

memblock_add(base, size);

```

⑦

③

在dtb中寻找带有"device_type" property的device node

④

如果在device node中找到了"device_type" property，则检查其property value是否为"memory"

这里③+④，就是为了定位

```
memory {

    device_type = "memory";

    reg = <0 0x00000000 0 0x40000000>; /* 1 GB */

};
```

这个device node。

⑤

获得device node的"reg" property并在下面的code中提取出base和size。

⑥

⑦

把从memory device node中提取的memory的base和size确定的region通过memblock_add()添加到"memory"的memblock_type中的region array中去。

=====

About memblock_reserve()

memblock_reserve()用于把不能用作allocate的memory区分出来 (reserved) 。

比如在arch/arm/mm/init.c中

```
void __init arm_memblock_init(const struct machine_desc *mdesc)
{
    /* Register the kernel text, kernel data and initrd with memblock. */

#ifdef CONFIG_XIP_KERNEL

    memblock_reserve(__pa(_sdata), _end - _sdata);

#else

    memblock_reserve(__pa(_stext), _end - _stext);
```

①

#endif

#ifdef CONFIG_BLK_DEV_INITRD

(A)

/* FDT scan will populate initrd_start */

if (initrd_start && !phys_initrd_size) {

 phys_initrd_start = __virt_to_phys(initrd_start);

 phys_initrd_size = initrd_end - initrd_start;

}

initrd_start = initrd_end = 0;

if (phys_initrd_size &&

 !memblock_is_region_memory(phys_initrd_start, phys_initrd_size)) {

 pr_err("INITRD: 0x%08llx+0x%08lx is not a memory region - disabling initrd\n",

 (u64)phys_initrd_start, phys_initrd_size);

 phys_initrd_start = phys_initrd_size = 0;

}

if (phys_initrd_size &&

 memblock_is_region_reserved(phys_initrd_start, phys_initrd_size)) {

 pr_err("INITRD: 0x%08llx+0x%08lx overlaps in-use memory region - disabling initrd\n",

 (u64)phys_initrd_start, phys_initrd_size);

 phys_initrd_start = phys_initrd_size = 0;

}

if (phys_initrd_size) {

 memblock_reserve(phys_initrd_start, phys_initrd_size);

/* Now convert initrd to virtual addresses */

initrd_start = __phys_to_virt(phys_initrd_start);

initrd_end = initrd_start + phys_initrd_size;

}

#endif

```
arm_mm_memblock_reserve();
```

 ②

```
/* reserve any platform specific memblock areas */
```

```
if (mdesc->reserve)
```

 ③

```
    mdesc->reserve();
```

```
early_init_fdt_scan_reserved_mem();
```

 ④

```
/*
```

```
 * reserve memory for DMA contiguous allocations,
```

```
 * must come from DMA area inside low memory
```

```
 */
```

```
dma_contiguous_reserve(arm_dma_limit);
```

```
arm_memblock_steal_permitted = false;
```

```
memblock_dump_all();
```

```
}
```

①

把kernel本身的code和readonly data部分所占的memory给reserve。

`_stext` and `_end` are in `vmlinux.lds`。

```
walterzh$ nm vmlinux-3.18.7-yocto-standard | grep "_end"
```

```
c06a68f4 B _end
```

```
walterzh$ nm vmlinux-3.18.7-yocto-standard | egrep _stext
```

```
c0008300 T _stext
```

```
memblock_reserve(__pa(_stext), _end - _stext) =
```

```
memblock_reserve(0xc0008300 - 0xc0000000, 0xc06a68f4 - 0xc0008300) =
```

memblock_reserve(0x8300, 0x69e5f4)

physical memory的[0x8300, 0x6a68f4)存放着kernel的code和readonly data.

(A)

如果由initramfs , 自然需要把initramfs所占的空间给reserved.

由于Granite2 and Gemstone2 LSP启动时的boot cmd为

```
mmc dev 1;ext2load mmc 1:2 0x400000 /boot/ulmage;ext2load mmc 1:2 0xf00000 /boot/mv6220-toc.dtb;setenv  
bootargs $bootargs root=/dev/mmcblk1p2 uio_pdrv_genirq.of_id=generic-uio rootwait;bootm 0x400000 - 0xf00000
```

这里bootm 0x400000 - 0xf00000中的"-"就代表initramfs。LSP没用到initramfs。

②

in arch/arm/mm/mmu.c

```
/*
```

```
 * Reserve the special regions of memory
```

```
*/
```

```
void __init arm_mm_memblock_reserve(void)
```

```
{
```

```
    /*
```

```
     * Reserve the page tables. These are already in use,
```

```
     * and can only be in node 0.
```

```
    */
```

```
    memblock_reserve(__pa(swapper_pg_dir), SWAPPER_PG_DIR_SIZE);
```

```
#ifdef CONFIG_SA1111
```

```

/*
 * Because of the SA1111 DMA bug, we want to preserve our
 * precious DMA-able memory...
 */

memblock_reserve(PHYS_OFFSET, __pa(swapper_pg_dir) - PHYS_OFFSET);

#endif

}

#ifdef CONFIG_ARM_LPAE

/* the first page is reserved for pgd */

#define SWAPPER_PG_DIR_SIZE    (PAGE_SIZE + \

                                PTRS_PER_PGD * PTRS_PER_PMD * sizeof(pmd_t))

#else

#define SWAPPER_PG_DIR_SIZE    (PTRS_PER_PGD * sizeof(pgd_t))

#endif

```

swapper_pg_dir是整个virtual memory的根，it is the virtual address of the initial page table.

相当于如下数组

```
pgd_t swapper_pg_dir[PTRS_PER_PGD];
```

```
#define PTRS_PER_PGD    2048
```

```
$ nm vmlinux-3.18.7-yocto-standard | grep swapper_pg_dir
```

```
c0003000 A swapper_pg_dir
```

```
__pa(swapper_pg_dir) = 0xc0003000 - 0xc0000000 = 0x00003000 (physical address)
```

```
SWAPPER_PG_DIR_SIZE = 2048 * 4 = 8K
```

即把swapper_pg_dir所占用的空间也reserve起来。

```
memblock_reserve(0x00003000 , 8K);
```

physical memory [0x3000, 0x5000)被swapper_pg_dir占用。

③

in arch/arm/mach-pegmatite/pegmatite.c

```
DT_MACHINE_START(PEGMATITE_DT, "Marvell Pegmatite (Device Tree)")
```

```
#ifdef CONFIG_SMP
```

```
    .smp          = smp_ops(pegmatite_smp_ops),
```

```
#endif
```

```
    .init_machine = pegmatite_dt_init,
```

```
    .map_io       = pegmatite_map_io,
```

```
    .init_early = pegmatite_init_early,
```

```
    .init_irq     = pegmatite_init_irq,
```

```
    .init_time    = pegmatite_timer_and_clk_init,
```

```
    .restart      = pegmatite_restart,
```

```
    .dt_compat    = pegmatite_dt_compat,
```

```
#ifdef CONFIG_ZONE_DMA
```

```
    .dma_zone_size = SZ_256M,
```

```
#endif
```

```
MACHINE_END
```

Gr2 LSP没有定义mdesc->reserve()的callback function。

④

in drivers/of/fdt.c

/**

* early_init_fdt_scan_reserved_mem() - create reserved memory regions

*

* This function grabs memory from early allocator for device exclusive use

* defined in device tree structures. It should be called by arch specific code

* once the early allocator (i.e. memblock) has been fully activated.

*/

void __init early_init_fdt_scan_reserved_mem(void)

{

int n;

u64 base, size;

if (!initial_boot_params)

return;

/* Reserve the dtb region */

early_init_dt_reserve_memory_arch(__pa(initial_boot_params), ⑤

fdt_totalsize(initial_boot_params),

0);

/* Process header /memreserve/ fields */

for (n = 0; ; n++) {

fdt_get_mem_rsv(initial_boot_params, n, &base, &size); ⑥

if (!size)

break;

early_init_dt_reserve_memory_arch(base, size, 0);

}

```

of_scan_flat_dt(__fdt_scan_reserved_mem, NULL);
fdt_init_reserved_mem();
}

```

⑦

⑧

⑤

dtb本身占用的memory要reserve。位于physical memory的0xf00000。

```
ext2load mmc 1:2 0xf00000 /boot/mv6220-toc.dtb;...;bootm 0x400000 - 0xf00000
```

⑥

dts中可以使用如下语法来reserve memory

```
/memreserve/ 0x20000000-0x21ffffff;
```

G2 LSP的dts没有用到。

⑦

在dtb中寻找reserved-memory device node。比如在mv6270-toc.dts中

```

reserved-memory {
    #address-cells = <2>;
    #size-cells = <2>;
    ranges;

    linux,cma {
        compatible = "shared-dma-pool";
        reusable;
        size = <0 0x10000>;
        alignment = <0 0x2000>;
    }
}

```

```

        linux,cma-default;

};

};

/**
 * fdt_scan_reserved_mem() - scan a single FDT node for reserved memory
 */
static int __init __fdt_scan_reserved_mem(unsigned long node, const char *uname,
                                           int depth, void *data)
{
    static int found;

    const char *status;

    int err;

    if (!found && depth == 1 && strcmp(uname, "reserved-memory") == 0) {      (A)
        if (__reserved_mem_check_root(node) != 0) {
            pr_err("Reserved memory: unsupported node format, ignoring\n");

            /* break scan */

            return 1;
        }

        found = 1;

        /* scan next node */

        return 0;
    } else if (!found) {

        /* scan next node */

        return 0;
    } else if (found && depth < 2) {

        /* scanning of /reserved-memory has been finished */

        return 1;
    }
}

```

```
}
```

```
status = of_get_flat_dt_prop(node, "status", NULL); (B)
```

```
if (status && strcmp(status, "okay") != 0 && strcmp(status, "ok") != 0)
```

```
    return 0;
```

```
err = __reserved_mem_reserve_reg(node, uname); (C)
```

```
if (err == -ENOENT && of_get_flat_dt_prop(node, "size", NULL))
```

```
    fdt_reserved_mem_save_node(node, uname, 0, 0); (D)
```

```
/* scan next node */
```

```
return 0;
```

```
}
```

```
CONFIG_OF_RESERVED_MEM=y
```

in drivers/base/dma-contiguous.c

```
obj-$(CONFIG_DMA_CMA) += dma-contiguous.o
```

CONFIG_DMA_CMA is not set.

in drivers/base/base/dma-coherent.c

```
obj-$(CONFIG_HAVE_GENERIC_DMA_COHERENT) += dma-coherent.o
```

```
CONFIG_HAVE_GENERIC_DMA_COHERENT=y
```

所以reserved-memory.linux,cma中的reserve memory由drivers/base/base/dma-coherent.c处理。

- (A)
- (B)
- (C)
- (D)

⑧

=====

kernel param "memblock=debug"可以打开memblock dump info.

```
# mmc dev 1;ext2load mmc 1:2 0x400000 /boot/ulimage;ext2load mmc 1:2 0xf00000 /boot/mv6220-toc.dtb;setenv  
bootargs $bootargs root=/dev/mmcblk1p2 uio_pdrv_genirq.of_id=generic-uio rootwait memblock=debug ;bootm  
0x400000 - 0xf00000
```

"memblock=debug" enable __memblock_dump_all()

```
1. memblock_reserve: [0x00000000008300-0x000000006a8af3] flags 0x0 arm_memblock_init+0x4c/0x1c0
2. memblock_reserve: [0x00000000003000-0x00000000007fff] flags 0x0 arm_memblock_init+0x174/0x1c0
3. memblock_reserve: [0x00000000f00000-0x00000000f11fff] flags 0x0 early_init_fdt_scan_reserved_me
  m+0x48/0x98
4. memblock_reserve: [0x00000000f00000-0x00000000f11fff] flags 0x0 early_init_fdt_scan_reserved_me
  m+0x78/0x98
5. memblock_reserve: [0x0000002f7f0000-0x0000002f7fffff] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
6. MEMBLOCK configuration:
7.   memory size = 0x3fc00000 reserved size = 0x6c77f4
8.   memory.cnt = 0x2
9.   memory[0x0][0x00000000000000-0x00000005ffffffff], 0x6000000 bytes flags: 0x0
10.  memory[0x1][0x000000006400000-0x00000003ffffffff], 0x39c00000 bytes flags: 0x0
11.  reserved.cnt = 0x4
12.  reserved[0x0][0x00000000003000-0x00000000007fff], 0x5000 bytes flags: 0x0
13.  reserved[0x1][0x00000000008300-0x000000006a8af3], 0x6a07f4 bytes flags: 0x0
14.  reserved[0x2][0x00000000f00000-0x00000000f11fff], 0x12000 bytes flags: 0x0
15.  reserved[0x3][0x0000002f7f0000-0x0000002f7fffff], 0x10000 bytes flags: 0x0
16. Forcing write-allocate cache policy for SMP
17. Memory policy: Data cache writealloc
18. memblock_reserve: [0x0000002f7ee000-0x0000002f7effff] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
19. memblock_reserve: [0x0000002f7ed000-0x0000002f7edfff] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
20. memblock_reserve: [0x0000002f7ecfa0-0x0000002f7ecfff] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
21. memblock_reserve: [0x0000002f7eb000-0x0000002f7ebfff] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
22. memblock_reserve: [0x0000002f7ea000-0x0000002f7eafff] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
23. memblock_reserve: [0x0000002f7ecf70-0x0000002f7ec9f] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
24. memblock_reserve: [0x0000002f7e9000-0x0000002f7e9fff] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
25. memblock_reserve: [0x0000002f7e8000-0x0000002f7e8fff] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
26. memblock_reserve: [0x0000002f7e7000-0x0000002f7e7fff] flags 0x0 memblock_alloc_range_nid+0x68/0
  x80
27. memblock_virt_alloc_try_nid_nopanic: 8388608 bytes align=0x0 nid=0 from=0x0 max_addr=0x0 alloc_
  node_mem_map.constprop.82+0x8c/0xb8
28. memblock_reserve: [0x0000002efe7000-0x0000002f7e6fff] flags 0x0 memblock_virt_alloc_internal+0x
```

16c/0x1a8

29. memblock_virt_alloc_try_nid_nopanic: 4 bytes align=0x0 nid=0 from=0x0 max_addr=0x0 free_area_in it_node+0x378/0x3fc
30. memblock_reserve: [0x0000002f7ecf40-0x0000002f7ecf43] flags 0x0 memblock_virt_alloc_internal+0x 16c/0x1a8
31. memblock_virt_alloc_try_nid_nopanic: 6144 bytes align=0x0 nid=0 from=0x0 max_addr=0x0 zone_wait _table_init+0x88/0xf4
32. memblock_reserve: [0x0000002efe5800-0x0000002efe6fff] flags 0x0 memblock_virt_alloc_internal+0x 16c/0x1a8
33. memblock_virt_alloc_try_nid_nopanic: 4 bytes align=0x0 nid=0 from=0x0 max_addr=0x0 free_area_in it_node+0x378/0x3fc
34. memblock_reserve: [0x0000002f7ecf00-0x0000002f7ecf03] flags 0x0 memblock_virt_alloc_internal+0x 16c/0x1a8
35. memblock_virt_alloc_try_nid_nopanic: 12288 bytes align=0x0 nid=0 from=0x0 max_addr=0x0 zone_wai t_table_init+0x88/0xf4
36. memblock_reserve: [0x0000002efe2800-0x0000002efe57ff] flags 0x0 memblock_virt_alloc_internal+0x 16c/0x1a8
37. memblock_virt_alloc_try_nid_nopanic: 4 bytes align=0x0 nid=0 from=0x0 max_addr=0x0 free_area_in it_node+0x378/0x3fc
38. memblock_reserve: [0x0000002f7ecfec0-0x0000002f7ecfec3] flags 0x0 memblock_virt_alloc_internal+0x 16c/0x1a8
39. memblock_virt_alloc_try_nid_nopanic: 12288 bytes align=0x0 nid=0 from=0x0 max_addr=0x0 zone_wai t_table_init+0x88/0xf4
40. memblock_reserve: [0x0000002efdf800-0x0000002efe27ff] flags 0x0 memblock_virt_alloc_internal+0x 16c/0x1a8
41. memblock_virt_alloc_try_nid: 40 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 setup_arch+0x5b0/0 xb14
42. memblock_reserve: [0x0000002f7ece80-0x0000002f7ecea7] flags 0x0 memblock_virt_alloc_internal+0x 16c/0x1a8
43. memblock_virt_alloc_try_nid: 40 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 setup_arch+0x5b0/0 xb14
44. memblock_reserve: [0x0000002f7ece40-0x0000002f7ece67] flags 0x0 memblock_virt_alloc_internal+0x 16c/0x1a8
45. memblock_reserve: [0x0000002efac5b0-0x0000002efdf7ff] flags 0x0 memblock_alloc_range_nid+0x68/0 x80
46. memblock_reserve: [0x0000002f7ecf50-0x0000002f7ecf6c] flags 0x0 memblock_alloc_range_nid+0x68/0 x80
47. memblock_reserve: [0x0000002f7ecf24-0x0000002f7ecf3e] flags 0x0 memblock_alloc_range_nid+0x68/0 x80
48. memblock_reserve: [0x0000002f7ecf08-0x0000002f7ecf22] flags 0x0 memblock_alloc_range_nid+0x68/0 x80
49. memblock_reserve: [0x0000002f7ecee4-0x0000002f7ecee] flags 0x0 memblock_alloc_range_nid+0x68/0 x80
50. memblock_reserve: [0x0000002f7ecfec8-0x0000002f7ecee2] flags 0x0 memblock_alloc_range_nid+0x68/0 x80

51. memblock_reserve: [0x0000002f7ecea8-0x0000002f7ecef] flags 0x0 memblock_alloc_range_nid+0x68/0x80

52. memblock_reserve: [0x0000002f7ecea8-0x0000002f7ecef] flags 0x0 memblock_alloc_range_nid+0x68/0x80

53. memblock_reserve: [0x0000002f7ecea8-0x0000002f7ecef] flags 0x0 memblock_alloc_range_nid+0x68/0x80

54. memblock_reserve: [0x0000002f7ecea8-0x0000002f7ecef] flags 0x0 memblock_alloc_range_nid+0x68/0x80

55. memblock_reserve: [0x0000002f7ecdf8-0x0000002f7ecef] flags 0x0 memblock_alloc_range_nid+0x68/0x80

56. memblock_reserve: [0x0000002f7ecde0-0x0000002f7ecdf7] flags 0x0 memblock_alloc_range_nid+0x68/0x80

57. memblock_reserve: [0x0000002f7ecdc8-0x0000002f7ecddf] flags 0x0 memblock_alloc_range_nid+0x68/0x80

58. memblock_reserve: [0x0000002f7ecdb0-0x0000002f7ecdc7] flags 0x0 memblock_alloc_range_nid+0x68/0x80

59. memblock_virt_alloc_try_nid: 133 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 start_kernel+0xd0/0x408

60. memblock_reserve: [0x0000002f7ecd00-0x0000002f7ecd84] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8

61. memblock_virt_alloc_try_nid: 133 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 start_kernel+0xfc/0x408

62. memblock_reserve: [0x0000002f7ecc40-0x0000002f7ecc4] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8

63. memblock_virt_alloc_try_nid: 133 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 start_kernel+0x128/0x408

64. memblock_reserve: [0x0000002f7ecb80-0x0000002f7ecc04] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8

65. memblock_virt_alloc_try_nid_nopanic: 4096 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 pcpu_all_oc_alloc_info+0x5c/0x98

66. memblock_reserve: [0x0000002efab580-0x0000002efac57f] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8

67. memblock_virt_alloc_try_nid_nopanic: 4096 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 pcpu_embed_first_chunk+0x4d8/0x768

68. memblock_reserve: [0x0000002efaa580-0x0000002efab57f] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8

69. memblock_virt_alloc_try_nid_nopanic: 81920 bytes align=0x1000 nid=-1 from=0x10000000 max_addr=0x0 pcpu_dfl_fc_alloc+0x6c/0x74

70. memblock_reserve: [0x0000002ef96000-0x0000002efa9fff] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8

71. __memblock_free_early: [0x0000002efa0000-0x0000002ef9ffff] pcpu_embed_first_chunk+0x618/0x768

72. __memblock_free_early: [0x0000002efaa000-0x0000002efa9fff] pcpu_embed_first_chunk+0x618/0x768

73. PERCPU: Embedded 10 pages/cpu @eef96000 s8384 r8192 d24384 u40960

74. memblock_virt_alloc_try_nid: 4 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 pcpu_setup_first_chunk+0x1c0/0x850


```

75. memblock_reserve: [0x0000002f7ecb40-0x0000002f7ecb43] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8
76. memblock_virt_alloc_try_nid: 4 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 pcpu_setup_first_chunk+0x1f0/0x850
77. memblock_reserve: [0x0000002f7ecb00-0x0000002f7ecb03] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8
78. memblock_virt_alloc_try_nid: 8 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 pcpu_setup_first_chunk+0x218/0x850
79. memblock_reserve: [0x0000002f7ecac0-0x0000002f7ecac7] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8
80. memblock_virt_alloc_try_nid: 8 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 pcpu_setup_first_chunk+0x240/0x850
81. memblock_reserve: [0x0000002f7eca80-0x0000002f7eca87] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8
82. memblock_virt_alloc_try_nid: 120 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 pcpu_setup_first_chunk+0x574/0x850
83. memblock_reserve: [0x0000002f7eca00-0x0000002f7eca77] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8
84. memblock_virt_alloc_try_nid: 68 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 pcpu_setup_first_chunk+0x5d0/0x850
85. memblock_reserve: [0x0000002f7ec980-0x0000002f7ec9c3] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8
86. memblock_virt_alloc_try_nid: 68 bytes align=0x0 nid=-1 from=0x0 max_addr=0x0 pcpu_setup_first_chunk+0x730/0x850
87. memblock_reserve: [0x0000002f7ec900-0x0000002f7ec943] flags 0x0 memblock_virt_alloc_internal+0x16c/0x1a8
88. __memblock_free_early: [0x0000002efab580-0x0000002efac57f] pcpu_embed_first_chunk+0x730/0x768
89. __memblock_free_early: [0x0000002efaa580-0x0000002efab57f] pcpu_embed_first_chunk+0x750/0x768

```

初始化期间那么多memblock_reserve(), 但size又很小, attachmen中有log.

```
root@granite2:~# ls -l /sys/kernel/debug/memblock
```

```
-r--r--r--  1 root  root      0 Jan  1  1970 memory
```

```
-r--r--r--  1 root  root      0 Jan  1  1970 reserved
```

```
root@granite2:~# cat /sys/kernel/debug/memblock/memory
```

```
0: 0x0000000000000000..0x0000000005ffffff      from 0 to 96M
```

```
1: 0x0000000006400000..0x000000003fffffff      from 100M to 1024M
```

也就是Linux看到的physical ram (address space)是断续的，分为两段。原因是from 96M to 100M是给R4用的。

```
root@granite2:~# cat /sys/kernel/debug/memblock/reserved
```

```
0: 0x00000000000003000..0x00000000000007fff      (swapper_pg_dir)
1: 0x00000000000008300..0x0000000000006a8af3      (kernel code and readonly data)
2: 0x00000000000f00000..0x00000000000f11fff      (dtb)
3: 0x0000000002eed2000..0x0000000002efa9fff
4: 0x0000000002efac5b0..0x0000000002f7ebfff
5: 0x0000000002f7ec900..0x0000000002f7ec943
6: 0x0000000002f7ec980..0x0000000002f7ec9c3
7: 0x0000000002f7eca00..0x0000000002f7eca77
8: 0x0000000002f7eca80..0x0000000002f7eca87
9: 0x0000000002f7ecac0..0x0000000002f7ecac7
10: 0x0000000002f7ecb00..0x0000000002f7ecb03
11: 0x0000000002f7ecb40..0x0000000002f7ecb43
12: 0x0000000002f7ecb80..0x0000000002f7ecc04
13: 0x0000000002f7ecc40..0x0000000002f7eccc4
14: 0x0000000002f7ecd00..0x0000000002f7ecd84
15: 0x0000000002f7ecdb0..0x0000000002f7ecec3
16: 0x0000000002f7ecec8..0x0000000002f7ecee2
17: 0x0000000002f7ecee4..0x0000000002f7ecefe
18: 0x0000000002f7ecf00..0x0000000002f7ecf03
19: 0x0000000002f7ecf08..0x0000000002f7ecf22
20: 0x0000000002f7ecf24..0x0000000002f7ecf3e
21: 0x0000000002f7ecf40..0x0000000002f7ecf43
22: 0x0000000002f7ecf50..0x0000000002f7ecf6c
23: 0x0000000002f7ecf70..0x0000000002f7fffff
```

start_kernel()

|

|

\ψ

setup_arch()

|

|

\ψ

arm_memblock_init()

in setup_arch() in arch/arm/kernel/setup.c

```

1. void __init setup_arch(char **cmdline_p)
2. {
3.     const struct machine_desc *mdesc;
4.
5.     setup_processor();
6.     mdesc = setup_machine_fdt(__atags_pointer);
7.     (1)
8.     if (!mdesc)
9.         mdesc = setup_machine_tags(__atags_pointer, __machine_arch_type); (2)
10.    machine_desc = mdesc;
11.    machine_name = mdesc->name;
12.
13.    if (mdesc->reboot_mode != REBOOT_HARD)
14.        reboot_mode = mdesc->reboot_mode;
15.
16.    init_mm.start_code = (unsigned long) _text;
17.    init_mm.end_code   = (unsigned long) _etext;
18.    init_mm.end_data   = (unsigned long) _edata;
19.    init_mm.brk        = (unsigned long) _end;
20.
21.    /* populate cmd_line too for later use, preserving boot_command_line */
22.    strcpy(cmd_line, boot_command_line, COMMAND_LINE_SIZE);
23.    *cmdline_p = cmd_line;
24.
25.    parse_early_param();
26.
27.    early_paging_init(mdesc, lookup_processor_type(read_cpuid_id())); (3)
28.    setup_dma_zone(mdesc);
29.    (4)
30.    sanity_check_meminfo();
31.    arm_memblock_init(mdesc); (5)
32.
33.    paging_init(mdesc);
34.
35.    .....

```

```
35. }
```

初始的关于memory information is from u-boot. (1)

在arm_memblock_init()中观察memblock的变化过程。(加入debug code)

```

1. void __init arm_memblock_init(const struct machine_desc *mdesc)
2. {
3.     /* Register the kernel text, kernel data and initrd with memblock. */
4.     printk("walter-0\n");
5.     __memblock_dump_all();
6.
7. #ifdef CONFIG_XIP_KERNEL
8.     memblock_reserve(__pa(_sdata), _end - _sdata);
9. #else
10.    memblock_reserve(__pa(_stext), _end - _stext);
11. #endif
12.
13.    printk("walter-1\n");
14.    __memblock_dump_all();
15.
16. #ifdef CONFIG_BLK_DEV_INITRD
17.     /* FDT scan will populate initrd_start */
18.     if (initrd_start && !phys_initrd_size) {
19.         phys_initrd_start = __virt_to_phys(initrd_start);
20.         phys_initrd_size = initrd_end - initrd_start;
21.     }
22.     initrd_start = initrd_end = 0;
23.     if (phys_initrd_size &&
24.         !memblock_is_region_memory(phys_initrd_start, phys_initrd_size)) {
25.         pr_err("INITRD: 0x%08llx+0x%08lx is not a memory region - disabling initrd\n",
26.             (u64)phys_initrd_start, phys_initrd_size);
27.         phys_initrd_start = phys_initrd_size = 0;
28.     }
29.     if (phys_initrd_size &&
30.         memblock_is_region_reserved(phys_initrd_start, phys_initrd_size)) {
31.         pr_err("INITRD: 0x%08llx+0x%08lx overlaps in-use memory region - disabling initrd\n",
32.             (u64)phys_initrd_start, phys_initrd_size);
33.         phys_initrd_start = phys_initrd_size = 0;
34.     }
35.     if (phys_initrd_size) {

```

```
36.         memblock_reserve(phys_initrd_start, phys_initrd_size);
37.
38.         /* Now convert initrd to virtual addresses */
39.         initrd_start = __phys_to_virt(phys_initrd_start);
40.         initrd_end = initrd_start + phys_initrd_size;
41.     }
42. #endif
43.     printk("walter-2\n");
44.     __memblock_dump_all();
45.
46.
47.     arm_mm_memblock_reserve();
48.
49.     printk("walter-3\n");
50.     __memblock_dump_all();
51.
52.
53.     /* reserve any platform specific memblock areas */
54.     if (mdesc->reserve)
55.         mdesc->reserve();
56.
57.     printk("walter-4\n");
58.     __memblock_dump_all();
59.
60.
61.     early_init_fdt_scan_reserved_mem();
62.
63.     printk("walter-5\n");
64.     __memblock_dump_all();
65.
66.
67.     /*
68.      * reserve memory for DMA contiguous allocations,
69.      * must come from DMA area inside low memory
70.      */
71.     dma_contiguous_reserve(arm_dma_limit);
```

```

72.
73.     printk("walter-6\n");
74.     __memblock_dump_all();
75.
76.     arm_memblock_steal_permitted = false;
77.     memblock_dump_all();
78. }

```

output如下(add comments) :

=====

walter-0

MEMBLOCK configuration:

memory size = 0x3fc00000 reserved size = 0x0

memory.cnt = 0x2

memory[0x0][0x0000000000000000-0x000000005fffffff], 0x6000000 bytes flags: 0x0

memory[0x1][0x00000006400000-0x00000003fffffff], 0x39c00000 bytes flags: 0x0

reserved.cnt = 0x1

reserved[0x0][0x0000000000000000-0xfffffffffffffff], 0x0 bytes flags: 0x0

Comments: 整个physical memory是1G - 4M (0x3fc00000),4M是被R4用的, 这块内存Linux
根本认为不存在。

1G - 4M的内存被分为2块。

from 0M to 96M

from 100M to 1G

no reserved spave

walter-1

MEMBLOCK configuration:

memory size = 0x3fc00000 reserved size = 0x6a07f4

memory.cnt = 0x2

memory[0x0][0x0000000000000000-0x000000005ffffff], 0x6000000 bytes flags: 0x0

memory[0x1][0x000000006400000-0x00000003ffffff], 0x39c00000 bytes flags: 0x0

reserved.cnt = 0x1

reserved[0x0][0x00000000008300-0x000000006a8af3], 0x6a07f4 bytes flags: 0x0

```
walterzh$ nm vmlinux-3.18.7-yocto-standard | grep _stext
```

```
c0008300 T _stext
```

```
walterzh$ nm vmlinux-3.18.7-yocto-standard | grep "_end$"
```

```
c06a8af4 B _end
```

symbol defines in vmlinux.lds

_stext 到_end包括了kernel的code和readonly data。它们被载入到virtual address space的[c0008300, c06a8af4) , 也就是physical address space的[0x00000000008300-0x000000006a8af3]。这段空间被reserved。不能被再利用。

walter-2

MEMBLOCK configuration:

memory size = 0x3fc00000 reserved size = 0x6a07f4

memory.cnt = 0x2

memory[0x0][0x0000000000000000-0x000000005ffffff], 0x6000000 bytes flags: 0x0

memory[0x1][0x000000006400000-0x00000003ffffff], 0x39c00000 bytes flags: 0x0

reserved.cnt = 0x1

reserved[0x0][0x00000000008300-0x000000006a8af3], 0x6a07f4 bytes flags: 0x0

没什么变化，因为虽然CONFIG_BLK_DEV_INITRD = y，但Gemstone2 LSP并没有initrd

walter-3

MEMBLOCK configuration:

memory size = 0x3fc00000 reserved size = 0x6a57f4

memory.cnt = 0x2

memory[0x0][0x0000000000000000-0x000000005ffffff], 0x6000000 bytes flags: 0x0

memory[0x1][0x00000006400000-0x00000003ffffff], 0x39c00000 bytes flags: 0x0

reserved.cnt = 0x2

reserved[0x0][0x00000000003000-0x00000000007fff], 0x5000 bytes flags: 0x0

reserved[0x1][0x00000000008300-0x000000006a8af3], 0x6a07f4 bytes flags: 0x0

```
1.  /*
2.   * Reserve the special regions of memory
3.   */
4.  void __init arm_mm_memblock_reserve(void)
5.  {
6.      /*
7.       * Reserve the page tables. These are already in use,
8.       * and can only be in node 0.
9.       */
10.     memblock_reserve(__pa(swapper_pg_dir), SWAPPER_PG_DIR_SIZE);
11.
12.     #ifdef CONFIG_SA1111
13.         /*
14.          * Because of the SA1111 DMA bug, we want to preserve our
15.          * precious DMA-able memory...
16.          */
17.         memblock_reserve(PHYS_OFFSET, __pa(swapper_pg_dir) - PHYS_OFFSET);
18.     #endif
19. }
```

walterzh\$ nm vmlinux-3.18.7-yocto-standard | grep swapper_pg_dir

c0003000 A swapper_pg_dir

```
#define SWAPPER_PG_DIR_SIZE    (PTRS_PER_PGD * sizeof(pgd_t))
```

```
#define PTRS_PER_PGD          2048
```

```
typedef struct { pmdval_t pgd[2]; } pgd_t;
```

==> SWAPPER_PG_DIR_SIZE = 2048 * 8 = 16K = 0x4000

这是first level page directory entry table的空间,同样被reserved。

Question : 真正的swapper_pg_dir应该开始于0xc0004000,占用[0xc0004000,0x8000),但reserved space是[0x3000, 0x8000)。

offset不对, size也多了1K。 Why ? ? ?

walter-4

MEMBLOCK configuration:

memory size = 0x3fc00000 reserved size = 0x6a57f4

memory.cnt = 0x2

memory[0x0][0x0000000000000000-0x000000005ffffff], 0x6000000 bytes flags: 0x0

memory[0x1][0x00000006400000-0x00000003ffffff], 0x39c00000 bytes flags: 0x0

reserved.cnt = 0x2

reserved[0x0][0x00000000003000-0x00000000007fff], 0x5000 bytes flags: 0x0

reserved[0x1][0x00000000008300-0x000000006a8af3], 0x6a07f4 bytes flags: 0x0

由于在machine descriptor中并没有customize .reserve() callback function,所以无变化。

walter-5

MEMBLOCK configuration:

memory size = 0x3fc00000 reserved size = 0x6c77f4

memory.cnt = 0x2

memory[0x0][0x0000000000000000-0x000000005ffffff], 0x6000000 bytes flags: 0x0

memory[0x1][0x000000006400000-0x00000003ffffff], 0x39c00000 bytes flags: 0x0

reserved.cnt = 0x4

reserved[0x0][0x00000000003000-0x0000000007fff], 0x5000 bytes flags: 0x0

reserved[0x1][0x00000000008300-0x000000006a8af3], 0x6a07f4 bytes flags: 0x0

reserved[0x2][0x00000000f00000-0x00000000f1ffff], 0x12000 bytes flags: 0x0

reserved[0x3][0x0000002f7f0000-0x0000002f7fffff], 0x10000 bytes flags: 0x0

```

1.  /**
2.   * early_init_fdt_scan_reserved_mem() - create reserved memory regions
3.   *
4.   * This function grabs memory from early allocator for device exclusive use
5.   * defined in device tree structures. It should be called by arch specific code
6.   * once the early allocator (i.e. memblock) has been fully activated.
7.   */
8.  void __init early_init_fdt_scan_reserved_mem(void)
9.  {
10.     int n;
11.     u64 base, size;
12.
13.     if (!initial_boot_params)
14.         return;
15.
16.     /* Reserve the dtb region */
17.     early_init_dt_reserve_memory_arch(__pa(initial_boot_params),
18.                                       fdt_totalsize(initial_boot_params),
19.                                       0);
20.
21.     /* Process header /memreserve/ fields */
22.     for (n = 0; ; n++) {
23.         fdt_get_mem_rsv(initial_boot_params, n, &base, &size);
24.         if (!size)
25.             break;
26.         early_init_dt_reserve_memory_arch(base, size, 0);
27.     }
28.
29.     of_scan_flat_dt(__fdt_scan_reserved_mem, NULL);
30.     fdt_init_reserved_mem();
31. }

```

/* Reserve the dtb region */

early_init_dt_reserve_memory_arch(__pa(initial_boot_params),

```
fdt_totalsize(initial_boot_params),
```

```
0);
```

reserve [0x00000000f00000-0x00000000f11fff] physical address space.

u-boot启动Linux kernel的命令如下：

```
bootcmd: mmc dev 1;ext2load mmc 1:2 0x400000 /boot/ulmage;ext2load mmc 1:2 0xf00000 /boot/mv6220-  
toc.dtb;setenv bootargs $bootargs root=/dev/mmcblk1p2 uio_pdrv_genirq.of_id=generic-uio rootwait;bootm  
0x400000 - 0xf00000
```

这里的0xf00000就是/boot/mv6220-toc.dtb的载入address。这里也要reserve。

```
reserved[0x3][0x0000002f7f0000-0x0000002f7fffff], 0x10000 bytes flags: 0x0
```

???

walter-6

MEMBLOCK configuration:

```
memory size = 0x3fc00000 reserved size = 0x6c77f4
```

```
memory.cnt = 0x2
```

```
memory[0x0][0x0000000000000000-0x000000005ffffff], 0x6000000 bytes flags: 0x0
```

```
memory[0x1][0x00000006400000-0x00000003ffffff], 0x39c00000 bytes flags: 0x0
```

```
reserved.cnt = 0x4
```

```
reserved[0x0][0x00000000003000-0x00000000007fff], 0x5000 bytes flags: 0x0
```

```
reserved[0x1][0x00000000008300-0x000000006a8af3], 0x6a07f4 bytes flags: 0x0
```

```
reserved[0x2][0x00000000f00000-0x00000000f11fff], 0x12000 bytes flags: 0x0
```

```
reserved[0x3][0x0000002f7f0000-0x0000002f7fffff], 0x10000 bytes flags: 0x0
```

`dma_contiguous_reserve()` function在Gemstone2 LSP上没有作用，memblock没有变化。

=====