

由于没有CONFIG_NO_BOOTMEM=y，所以在memblock与buddy allocator之间没有boot memory manager的参与。

start_kernel()

|

|

\\

mm_init()

|

|

\\

mem_init()

在mem_init()中完成内存管理从memblock到buddy allocator的过渡。

在mem_init()以前，所有memory的状况都记录在struct memblock **memblock**这个global variable中

in mm/memblock.c

```
struct memblock memblock __initdata_memblock = {  
    .memory.regions      = memblock_memory_init_regions,  
    .memory.cnt          = 1, /* empty dummy entry */  
    .memory.max          = INIT_MEMBLOCK_REGIONS,  
  
    .reserved.regions    = memblock_reserved_init_regions,
```

```

.reserved.cnt      = 1, /* empty dummy entry */

.reserved.max      = INIT_MEMBLOCK_REGIONS,

#ifdef CONFIG_HAVE_MEMBLOCK_PHYS_MAP

.physmem.regions   = memblock_physmem_init_regions,

.physmem.cnt       = 1, /* empty dummy entry */

.physmem.max       = INIT_PHYSMEM_REGIONS,

#endif

.bottom_up         = false,

.current_limit     = MEMBLOCK_ALLOC_ANYWHERE,

};

```

1. 所有可利用的memory都在memblock.memory.regions[]这个数组中，memblock.memory.cnt是这个array中的有效项数。每一项都是一个memblock_region。
memblock.memory.regions[memblock.memory.cnt]

```

struct memblock_region {

    phys_addr_t base;

    phys_addr_t size;

    unsigned long flags;

#ifdef CONFIG_HAVE_MEMBLOCK_NODE_MAP

    int nid;

#endif

};

```

base and size指定了一块region(physical memory).

并且这些entry都是按memblock_region的base排序的（从小到大）。

2. 所有已经被reserved,即不能作为free memory分配的空间都记录

在**memblock.reserved.regions[]**这个数组中，同样memblock.reserved.cnt是这个array的有效项数。这些entry同样都是memblock_region的base排序的（从小到大），**memblock.reserved.regions[memblock.reserved.cnt]**

也就是说从memblock.memory.regions[]所表示的memory中去除memblock.reserved.regions[]记录的reserved空间就是可以用allocate的free memory!

dump memblock status info before mem_init()

MEMBLOCK configuration:

memory size = 0x3fc00000 reserved size = 0xfe0d3b

memory.cnt = 0x2

memory[0x0][0x0000000000000000-0x000000005ffffff], 0x6000000 bytes flags: 0x0

memory[0x1][0x00000006400000-0x00000003ffffff], 0x39c00000 bytes flags: 0x0

reserved.cnt = 0x18

reserved[0x0][0x00000000003000-0x00000000007fff], 0x5000 bytes flags: 0x0

reserved[0x1][0x00000000008300-0x0000000006a68f3], 0x69e5f4 bytes flags: 0x0

reserved[0x2][0x00000000f00000-0x00000000f11fff], 0x12000 bytes flags: 0x0

reserved[0x3][0x0000002eed1000-0x0000002efa8fff], 0xd8000 bytes flags: 0x0

reserved[0x4][0x0000002efabd68-0x0000002f7ebfff], 0x840298 bytes flags: 0x0

reserved[0x5][0x0000002f7ec9c0-0x0000002f7eca03], 0x44 bytes flags: 0x0

reserved[0x6][0x0000002f7eca40-0x0000002f7eca83], 0x44 bytes flags: 0x0

reserved[0x7][0x0000002f7ecac0-0x0000002f7ecb37], 0x78 bytes flags: 0x0

reserved[0x8][0x0000002f7ecb40-0x0000002f7ecb47], 0x8 bytes flags: 0x0

reserved[0x9][0x0000002f7ecb80-0x0000002f7ecb87], 0x8 bytes flags: 0x0
reserved[0xa][0x0000002f7ecbc0-0x0000002f7ecbc3], 0x4 bytes flags: 0x0
reserved[0xb][0x0000002f7ecc00-0x0000002f7ecc75], 0x76 bytes flags: 0x0
reserved[0xc][0x0000002f7ecc80-0x0000002f7eccf5], 0x76 bytes flags: 0x0
reserved[0xd][0x0000002f7ecd00-0x0000002f7ecd75], 0x76 bytes flags: 0x0
reserved[0xe][0x0000002f7ecd80-0x0000002f7ecd83], 0x4 bytes flags: 0x0
reserved[0xf][0x0000002f7ecdb0-0x0000002f7ecec3], 0x114 bytes flags: 0x0
reserved[0x10][0x0000002f7ecec8-0x0000002f7ecee2], 0x1b bytes flags: 0x0
reserved[0x11][0x0000002f7ecee4-0x0000002f7ecefe], 0x1b bytes flags: 0x0
reserved[0x12][0x0000002f7ecf00-0x0000002f7ecf03], 0x4 bytes flags: 0x0
reserved[0x13][0x0000002f7ecf08-0x0000002f7ecf22], 0x1b bytes flags: 0x0
reserved[0x14][0x0000002f7ecf24-0x0000002f7ecf3e], 0x1b bytes flags: 0x0
reserved[0x15][0x0000002f7ecf40-0x0000002f7ecf43], 0x4 bytes flags: 0x0
reserved[0x16][0x0000002f7ecf50-0x0000002f7ecf6c], 0x1d bytes flags: 0x0
reserved[0x17][0x0000002f7ecf70-0x0000002f7fffff], 0x13090 bytes flags: 0x0

in arch/arm/mm/init.c

/*

* mem_init() marks the free areas in the mem_map and tells us how much

* memory is free. This is done after various parts of the system have

* claimed their memory after the kernel image.

*/

void __init mem_init(void)

{

#ifdef CONFIG_HAVE_TCM

/* These pointers are filled in on TCM detection */

extern u32 dtcm_end;

```
extern u32 itcm_end;
```

```
#endif
```

```
set_max_mapnr(pfn_to_page(max_pfn) - mem_map); ①
```

```
/* this will put all unused low memory onto the freelists */
```

```
free_unused_memmap(); ②
```

```
free_all_bootmem(); ③
```

```
#ifdef CONFIG_SA1111
```

```
/* now that our DMA memory is actually so designated, we can free it */
```

```
free_reserved_area(__va(PHYS_OFFSET), swapper_pg_dir, -1, NULL);
```

```
#endif
```

```
free_highpages(); ④
```

```
mem_init_print_info(NULL); ⑤
```

```
#define MLK(b, t) b, t, ((t) - (b)) >> 10 ⑥
```

```
#define MLM(b, t) b, t, ((t) - (b)) >> 20
```

```
#define MLK_ROUNDUP(b, t) b, t, DIV_ROUND_UP(((t) - (b)), SZ_1K)
```

```
printk(KERN_NOTICE "Virtual kernel memory layout:\n"
```

```
    "    vector : 0x%08lx - 0x%08lx  (%4ld kB)\n"
```

```
#ifdef CONFIG_HAVE_TCM
```

" DTCM : 0x%08lx - 0x%08lx (%4ld kB)\n"

" ITCM : 0x%08lx - 0x%08lx (%4ld kB)\n"

#endif

" fixmap : 0x%08lx - 0x%08lx (%4ld kB)\n"

" vmalloc : 0x%08lx - 0x%08lx (%4ld MB)\n"

" lowmem : 0x%08lx - 0x%08lx (%4ld MB)\n"

#ifdef CONFIG_HIGHMEM

" pkmap : 0x%08lx - 0x%08lx (%4ld MB)\n"

#endif

#ifdef CONFIG_MODULES

" modules : 0x%08lx - 0x%08lx (%4ld MB)\n"

#endif

" .text : 0x%p" " - 0x%p" " (%4td kB)\n"

" .init : 0x%p" " - 0x%p" " (%4td kB)\n"

" .data : 0x%p" " - 0x%p" " (%4td kB)\n"

" .bss : 0x%p" " - 0x%p" " (%4td kB)\n",

MLK(UL(CONFIG_VECTORS_BASE), UL(CONFIG_VECTORS_BASE) +
(PAGE_SIZE)),

#ifdef CONFIG_HAVE_TCM

MLK(DTCM_OFFSET, (unsigned long) dtcm_end),

MLK(ITCM_OFFSET, (unsigned long) itcm_end),

#endif

MLK(FIXADDR_START, FIXADDR_TOP),

MLM(VMALLOC_START, VMALLOC_END),

```

        MLM(PAGE_OFFSET, (unsigned long)high_memory),

#ifdef CONFIG_HIGHMEM

        MLM(PKMAP_BASE, (PKMAP_BASE) + (LAST_PKMAP) *

            (PAGE_SIZE)),

#endif

#ifdef CONFIG_MODULES

        MLM(MODULES_VADDR, MODULES_END),

#endif

        MLK_ROUNDUP(_text, _etext),

        MLK_ROUNDUP(__init_begin, __init_end),

        MLK_ROUNDUP(_sdata, _edata),

        MLK_ROUNDUP(__bss_start, __bss_stop));

#undef MLK

#undef MLM

#undef MLK_ROUNDUP

/*

 * Check boundaries twice: Some fundamental inconsistencies can

 * be detected at build time already.

 */

#ifdef CONFIG_MMU

        BUILD_BUG_ON(TASK_SIZE > MODULES_VADDR);

        BUG_ON(TASK_SIZE > MODULES_VADDR);

```

```
#endif
```

```
#ifdef CONFIG_HIGHMEM
```

```
    BUILD_BUG_ON(PKMAP_BASE + LAST_PKMAP * PAGE_SIZE > PAGE_OFFSET);
```

```
    BUG_ON(PKMAP_BASE + LAST_PKMAP * PAGE_SIZE > PAGE_OFFSET);
```

```
#endif
```

```
if (PAGE_SIZE >= 16384 && get_num_physpages() <= 128) {
```

```
    extern int sysctl_overcommit_memory;
```

```
    /*
```

```
    * On a machine this small we won't get
```

```
    * anywhere without overcommit, so turn
```

```
    * it on by default.
```

```
    */
```

```
    sysctl_overcommit_memory = OVERCOMMIT_ALWAYS;
```

```
}
```

```
}
```

①

设置mem_map[] array的最大有效index，即max_mapnr。

```
unsigned long max_mapnr;
```

max_pfn是最大的page frame,而pfn_to_page(max_pfn)转换page frame到对应的struct page。


```
struct page *mem_map;
```

one page frame, one struct page.

mem_map[]是struct page的array。

```
438e08f1-r0/image/boot$ nm vmlinux-3.18.7-yocto-standard | grep " mem_map"
```

```
c0667080 B mem_map
```

```
mem_map[0]    -->    physical page 0
```

```
mem_map[1]    --> physical page 1
```

```
.....
```

```
mem_map[262143] --> physical page 262143 (Gr2有1G memory)
```

```
max_mapnr = 262144 ;
```

```
index of mem_map[] < max_mapnr
```

```
②
```

```
/*
```

```
* The mem_map array can get very big. Free the unused area of the memory map.
```

```
*/
```

```
static void __init free_unused_memmap(void)
```

```

{

unsigned long start, prev_end = 0;

struct memblock_region *reg;


/*

* This relies on each bank being in address order.

* The banks are sorted previously in bootmem_init().

*/

for_each_memblock(memory, reg) {                                (A)

    start = memblock_region_memory_base_pfn(reg);                (A.1)


#ifdef CONFIG_SPARSEMEM

/*

* Take care not to free memmap entries that don't exist

* due to SPARSEMEM sections which aren't present.

*/

start = min(start,

                ALIGN(prev_end, PAGE_SIZE));

#else

/*

* Align down here since the VM subsystem insists that the

* memmap entries are valid from the bank start aligned to

* MAX_ORDER_NR_PAGES.

*/

start = round_down(start, MAX_ORDER_NR_PAGES);                (B)

```

```
#endif
```

```
/*
```

```
 * If we had a previous bank, and there is a space
```

```
 * between the current bank and the previous, free it.
```

```
*/
```

```
if (prev_end && prev_end < start)                (C)
```

```
    free_memmap(prev_end, start);
```

```
/*
```

```
 * Align up here since the VM subsystem insists that the
```

```
 * memmap entries are valid from the bank end aligned to
```

```
 * MAX_ORDER_NR_PAGES.
```

```
*/
```

```
prev_end = ALIGN(memblock_region_memory_end_pfn(reg), (D)
```

```
    MAX_ORDER_NR_PAGES);
```

```
}
```

```
#ifdef CONFIG_SPARSEMEM
```

```
    if (!IS_ALIGNED(prev_end, PAGE_SIZE))
```

```
        free_memmap(prev_end,
```

```
            ALIGN(prev_end, PAGE_SIZE));
```

```
#endif
```

```
}
```

(A)

对memblock.memory.regions[] array进行enumerate.

memory size = 0x3fc00000 reserved size = 0xfe0d3b

memory.cnt = 0x2

memory[0x0][0x0000000000000000-0x000000005ffffff], 0x6000000 bytes flags: 0x0

memory[0x1][0x00000006400000-0x00000003ffffff], 0x39c00000 bytes flags: 0x0

memblock.memory.regions[0].base = 0

memblock.memory.regions[0].size = 0x6000000 (96M)

memblock.memory.regions[0].base = 0x6400000

memblock.memory.regions[0].size = 0x39c00000 (924M)

当中没挖去的4M给Secure core R4用的。

(A.1)

使得base向上取整到page边界。比如base = 5000,向上取整后就是8192 (0x2000)

对memblock.memory.regions[0].base = 0而言, start = 0

对memblock.memory.regions[1].base = 0x6400000 (100M)而言, start = 0x6400000(本来就在page alignment).

(B)

CONFIG_FORCE_MAX_ZONEORDER=16

in include/linux/mmzone.h

```
/* Free memory management - zoned buddy allocator. */
```

```
#ifndef CONFIG_FORCE_MAX_ZONEORDER
```

```
#define MAX_ORDER 11
```

```
#else
```

```
#define MAX_ORDER CONFIG_FORCE_MAX_ZONEORDER
```

```
#endif
```

```
#define MAX_ORDER_NR_PAGES (1 << (MAX_ORDER - 1))
```

MAX_ORDER_NR_PAGES为15, 即 2^{15} pages = 32K pages = **128M**

再向下取整到32K page边界, 即128M的边界。

对start = 0而言, start取整后还是0

对start = 0x6400000而言, $100M < 128M \Rightarrow start = 0$

??? 是否代表在大可分配连续physical memory为128M呢 ???

??? Linux kernel可以分配的最大连续physical memory是多大 ???

(C)

在enumerate memblock.memory.regions[0]时, prev_end = 0

在enumerate memblock.memory.regions[1]时, prev_end = 0

所以条件都不满足。也就是free_unused_memmap()什么都没做, 等于没调用一样!

(D)

对region上界(base + size)在128M上向下取整。

由于memblock.memory.regions[0].base + memblock.memory.regions[0].size = 96M < 128M

==>

prev_end = 0

③

真正完成从memblock到buddy allocator的交接在free_all_bootmem() function中。

in mm/nobootmem.c

```
/**
```

```
 * free_all_bootmem - release free pages to the buddy allocator
```

```
 *
```

```
 * Returns the number of pages actually released.
```

```
 */
```

```
unsigned long __init free_all_bootmem(void)
```

```
{
```

```
    unsigned long pages;
```

```
    reset_all_zones_managed_pages();
```

```
    /*
```

```
     * We need to use NUMA_NO_NODE instead of NODE_DATA(0)->node_id
```

* because in some case like Node0 doesn't have RAM installed

* low ram will be on Node1

*/

pages = free_low_memory_core_early();

totalram_pages += pages;

return pages;

}

in mm/nobootmem.c

static unsigned long __init free_low_memory_core_early(void)

{

unsigned long count = 0;

phys_addr_t start, end;

u64 i;

memblock_clear_hotplug(0, -1);

for_each_free_mem_range(i, NUMA_NO_NODE, &start, &end, NULL) (A)

count += __free_memory_core(start, end); (B)

#ifdef CONFIG_ARCH_DISCARD_MEMBLOCK

{

phys_addr_t size;

```

/* Free memblock.reserved array if it was allocated */

size = get_allocated_memblock_reserved_regions_info(&start);

if (size)

    count += __free_memory_core(start, start + size);


/* Free memblock.memory array if it was allocated */

size = get_allocated_memblock_memory_regions_info(&start);

if (size)

    count += __free_memory_core(start, start + size);

}

#endif


return count;

}

```

(A)

(B)

/**

* for_each_free_mem_range - iterate through free memblock areas

* @i: u64 used as loop variable

* @nid: node selector, %NUMA_NO_NODE for all nodes

* @p_start: ptr to phys_addr_t for start address of the range, can be %NULL

* @p_end: ptr to phys_addr_t for end address of the range, can be %NULL


```

* @p_nid: ptr to int for nid of the range, can be %NULL
*
* Walks over free (memory && !reserved) areas of memblock. Available as
* soon as memblock is initialized.
*/

#define for_each_free_mem_range(i, nid, p_start, p_end, p_nid) \
    for_each_mem_range(i, &memblock.memory, &memblock.reserved, \
        nid, p_start, p_end, p_nid)

==>

/**
 * for_each_mem_range - iterate through memblock areas from type_a and not
 * included in type_b. Or just type_a if type_b is NULL.
 *
 * @i: u64 used as loop variable
 *
 * @type_a: ptr to memblock_type to iterate
 *
 * @type_b: ptr to memblock_type which excludes from the iteration
 *
 * @nid: node selector, %NUMA_NO_NODE for all nodes
 *
 * @p_start: ptr to phys_addr_t for start address of the range, can be %NULL
 *
 * @p_end: ptr to phys_addr_t for end address of the range, can be %NULL
 *
 * @p_nid: ptr to int for nid of the range, can be %NULL
 */

#define for_each_mem_range(i, type_a, type_b, nid, \
    p_start, p_end, p_nid) \

```

```

for (i = 0, __next_mem_range(&i, nid, type_a, type_b, \
                             p_start, p_end, p_nid); \
     i != (u64)ULLONG_MAX; \
     __next_mem_range(&i, nid, type_a, type_b, \
                     p_start, p_end, p_nid))

```

in mm/memblock.c

/**

* __next_mem_range_rev - generic next function for for_each_*_range_rev()

*

* Finds the next range from type_a which is not marked as unsuitable

* in type_b.

*

* @idx: pointer to u64 loop variable

* @nid: nid: node selector, %NUMA_NO_NODE for all nodes

* @type_a: pointer to memblock_type from where the range is taken

* @type_b: pointer to memblock_type which excludes memory from being taken

* @out_start: ptr to phys_addr_t for start address of the range, can be %NULL

* @out_end: ptr to phys_addr_t for end address of the range, can be %NULL

* @out_nid: ptr to int for nid of the range, can be %NULL

*

* Reverse of __next_mem_range().

*/

```

void __init_memblock __next_mem_range_rev(u64 *idx, int nid,

        struct memblock_type *type_a,

        struct memblock_type *type_b,

        phys_addr_t *out_start,

        phys_addr_t *out_end, int *out_nid)

```

???

④

```

static void __init free_highpages(void)

{

#ifdef CONFIG_HIGHMEM

    unsigned long max_low = max_low_pfn;

    struct memblock_region *mem, *res;

    /* set highmem page free */

    for_each_memblock(memory, mem) {

        unsigned long start = memblock_region_memory_base_pfn(mem);

        unsigned long end = memblock_region_memory_end_pfn(mem);

        /* Ignore complete lowmem entries */

        if (end <= max_low)

            continue;

```

(A)

```
/* Truncate partial highmem entries */
```

```
if (start < max_low)
```

```
    start = max_low;
```

(B)

```
/* Find and exclude any reserved regions */
```

```
for_each_memblock(reserved, res) {
```

(C)

```
    unsigned long res_start, res_end;
```

```
    res_start = memblock_region_reserved_base_pfn(res);
```

(D)

```
    res_end = memblock_region_reserved_end_pfn(res);
```

(E)

```
    if (res_end < start)
```

(F)

```
        continue;
```

```
    if (res_start < start)
```

(G)

```
        res_start = start;
```

```
    if (res_start > end)
```

(H)

```
        res_start = end;
```

```
    if (res_end > end)
```

(I)

```
        res_end = end;
```

```
    if (res_start != start)
```

```
        free_area_high(start, res_start);
```

```
    start = res_end;
```

```
    if (start == end)
```

```

        break;

    }

    /* And now free anything which remains */

    if (start < end)

        free_area_high(start, end);

    }

#endif

}

```

(A)

对memblock.memory.regions[] array enumerate。

memblock.memory.regions[0]不满足条件。

(B)

start = memblock.memory.regions[1].base = 100M < max_low,so start = max_low

(C)

在memblock.memory.regions[1]中的memory从max_low到(memblock.memory.regions[1].base + memblock.memory.regions[1].size)属于high memory , 但要去出reserved regions,即 memblock.reserved.regions[] array中的space。

reserved.cnt = 0x18

reserved[0x0][0x000000000003000-0x000000000007fff], 0x5000 bytes flags: 0x0

reserved[0x1][0x00000000008300-0x000000006a68f3], 0x69e5f4 bytes flags: 0x0

reserved[0x2][0x00000000f00000-0x00000000f11fff], 0x12000 bytes flags: 0x0

下面的reserved space位于high memory

reserved[0x3][0x0000002eed1000-0x0000002efa8fff], 0xd8000 bytes flags: 0x0

reserved[0x4][0x0000002efabd68-0x0000002f7ebfff], 0x840298 bytes flags: 0x0

reserved[0x5][0x0000002f7ec9c0-0x0000002f7eca03], 0x44 bytes flags: 0x0

reserved[0x6][0x0000002f7eca40-0x0000002f7eca83], 0x44 bytes flags: 0x0

reserved[0x7][0x0000002f7ecac0-0x0000002f7ecb37], 0x78 bytes flags: 0x0

reserved[0x8][0x0000002f7ecb40-0x0000002f7ecb47], 0x8 bytes flags: 0x0

reserved[0x9][0x0000002f7ecb80-0x0000002f7ecb87], 0x8 bytes flags: 0x0

reserved[0xa][0x0000002f7ecbc0-0x0000002f7ecbc3], 0x4 bytes flags: 0x0

reserved[0xb][0x0000002f7ecc00-0x0000002f7ecc75], 0x76 bytes flags: 0x0

reserved[0xc][0x0000002f7ecc80-0x0000002f7eccf5], 0x76 bytes flags: 0x0

reserved[0xd][0x0000002f7ecd00-0x0000002f7ecd75], 0x76 bytes flags: 0x0

reserved[0xe][0x0000002f7ecd80-0x0000002f7ecd83], 0x4 bytes flags: 0x0

reserved[0xf][0x0000002f7ecdb0-0x0000002f7ecec3], 0x114 bytes flags: 0x0

reserved[0x10][0x0000002f7ecec8-0x0000002f7ecee2], 0x1b bytes flags: 0x0

reserved[0x11][0x0000002f7ecee4-0x0000002f7ecefe], 0x1b bytes flags: 0x0

reserved[0x12][0x0000002f7ecf00-0x0000002f7ecf03], 0x4 bytes flags: 0x0

reserved[0x13][0x0000002f7ecf08-0x0000002f7ecf22], 0x1b bytes flags: 0x0

reserved[0x14][0x0000002f7ecf24-0x0000002f7ecf3e], 0x1b bytes flags: 0x0

reserved[0x15][0x0000002f7ecf40-0x0000002f7ecf43], 0x4 bytes flags: 0x0

reserved[0x16][0x0000002f7ecf50-0x0000002f7ecf6c], 0x1d bytes flags: 0x0

reserved[0x17][0x0000002f7ecf70-0x0000002f7fffff], 0x13090 bytes flags: 0x0

对位于high memory空间的reserved space进行enumerate。

(D)

(E)

res_start, res_end都对齐在page边界上，并且包括了reserved的region。

(F)

表示该reserved region不再high memory之内

(G)

表示该reserved region的下边界（低地址的一端）落在high memory的外面（即在normal memory）

(H)

表示该reserved region的下边界（低地址的一端）超出了当前memory bank的end边界（reserved region跨memory bank，应该不大可能）

也就是reserved region根本不在当前memory bank之内

(I)

表示该reserved region的上边界（高地址的一端）超出了当前memory bank的end边界

code看上去由很多判断，实际在G2 LSP上没那么复杂，就是从memblock.memory.regions[1]代表的内存中从high memory的边界开始，凡是没有reserved region的身影的page都调用free_area_high()。

```
#ifdef CONFIG_HIGHMEM
```

```
static inline void free_area_high(unsigned long pfn, unsigned long end)
```

```
{
```



```

    for (; pfn < end; pfn++)

        free_highmem_page(pfn_to_page(pfn));

}

#endif

#ifdef CONFIG_HIGHMEM

void free_highmem_page(struct page *page)

{

    __free_reserved_page(page);

    totalram_pages++;

    page_zone(page)->managed_pages++;

    totalhigh_pages++;

}

#endif

```

⑤

该函数print如下信息：

Memory: 1028212K/1044480K available (4550K kernel code, 217K rwdatas, 1352K rodata, 208K init, 448K bss, 16268K reserved, 270336K highmem)

```

void __init mem_init_print_info(const char *str)

{

    unsigned long physpages, codesize, datasize, rosize, bss_size;

    unsigned long init_code_size, init_data_size;

```

```

physpages = get_num_physpages();

codesize = _etext - _stext;

datasize = _edata - _sdata;

rosize = __end_rodata - __start_rodata;

bss_size = __bss_stop - __bss_start;

init_data_size = __init_end - __init_begin;

init_code_size = _einittext - _sinittext;

```

```

/*

```

```

 * Detect special cases and adjust section sizes accordingly:

```

```

 * 1) .init.* may be embedded into .data sections

```

```

 * 2) .init.text.* may be out of [__init_begin, __init_end],

```

```

 *   please refer to arch/tile/kernel/vmlinux.lds.S.

```

```

 * 3) .rodata.* may be embedded into .text or .data sections.

```

```

 */

```

```

#define adj_init_size(start, end, size, pos, adj) \

```

```

do { \

```

```

    if (start <= pos && pos < end && size > adj) \

```

```

        size -= adj; \

```

```

} while (0)

```

```

adj_init_size(__init_begin, __init_end, init_data_size,

```

(A)

```

    _sinittext, init_code_size);

```

```

adj_init_size(_stext, _etext, codesize, _sinittext, init_code_size); (B)

```

```
adj_init_size(_sdata, _edata, datasize, __init_begin, init_data_size);(C)
```

```
adj_init_size(_stext, _etext, codesize, __start_rodata, rosize);          (D)
```

```
adj_init_size(_sdata, _edata, datasize, __start_rodata, rosize);          (E)
```

```
#undef  adj_init_size
```

```
printf("Memory: %luK/%luK available "
```

```
      "(%luK kernel code, %luK rdata, %luK rodata, "
```

```
      "%luK init, %luK bss, %luK reserved"
```

```
#ifdef  CONFIG_HIGHMEM
```

```
      ", %luK highmem"
```

```
#endif
```

```
      "%s%s)\n",
```

```
      nr_free_pages() << (PAGE_SHIFT-10), physpages << (PAGE_SHIFT-10),
```

```
      codesize >> 10, datasize >> 10, rosize >> 10,
```

```
      (init_data_size + init_code_size) >> 10, bss_size >> 10,
```

```
      (physpages - totalram_pages) << (PAGE_SHIFT-10),
```

```
#ifdef  CONFIG_HIGHMEM
```

```
      totalhigh_pages << (PAGE_SHIFT-10),
```

```
#endif
```

```
      str ? ", " : "", str ? str : "");
```

```
}
```

(A)

```
if(__init_begin < _sinittext && _sinittext < __init_end && init_data_size > init_code_size)
```

```
init_data_size -= init_code_size;
```

表示_sinittext是被包含在(__init_begin, __init_end)之间的，自然init_data_size应该减去init_code_size的大小（否则就重复计算了）

(B)

如果_sinittext被包含在_s.text里，则也要调整大小。

(C)

如果__init_begin被包含在_s.data里，则也要调整大小。

(D)

如果__start_rodata被包含在_s.text里，则也要调整大小。

(E)

如果__start_rodata被包含在_s.data里，则也要调整大小。

physpages ???

totalram_pages ???

totalhigh_pages ???

in Gr2 LSP, memory info as follow

Memory: 1028212K/1044480K available (4550K kernel code, 217K rwd data, 1352K rodata, 208K init, 448K bss, 16268K reserved, 270336K highmem)

Virtual kernel memory layout:

vector : 0xffff0000 - 0xffff1000 (4 kB)

fixmap : 0xffc00000 - 0xffe00000 (2048 kB)

vmalloc : 0xf0000000 - 0xff000000 (240 MB)

lowmem : 0xc0000000 - 0xef800000 (760 MB)

pkmap : 0xbfe00000 - 0xc0000000 (2 MB)

modules : 0xbf000000 - 0xbfe00000 (14 MB)

.text : 0xc0008000 - 0xc05cbdd8 (5904 kB)

.init : 0xc05cc000 - 0xc0600000 (208 kB)

.data : 0xc0600000 - 0xc0636518 (218 kB)

.bss : 0xc0636518 - 0xc06a68f4 (449 kB)