

in arch/arm/kernel/setup.c

```
1. void __init setup_arch(char **cmdline_p)
2. {
3.     .....
4.
5.     unflatten_device_tree();
6.
7.     arm_dt_init_cpu_maps();
8.     psci_init();
9. #ifdef CONFIG_SMP
10.     if (is_smp()) {
11.         if (!mdesc->smp_init || !mdesc->smp_init()) {
12.             if (psci_smp_available())
13.                 smp_set_ops(&psci_smp_ops);
14.             else if (mdesc->smp)
15.                 smp_set_ops(mdesc->smp);
16.         }
17.         smp_init_cpus();
18.         smp_build_mpidr_hash();
19.     }
20. #endif
21.
22.     .....
23. }
```

在G2 LSP中mdesc->smp\_init没有定义。

```
1. DT_MACHINE_START(PEGMATITE_DT, "Marvell Pegmatite (Device Tree)")
2.     .dt_compat      = pegmatite_dt_compat,
3. #ifdef CONFIG_ZONE_DMA
4.     .dma_zone_size  = SZ_256M,
5. #endif
6. MACHINE_END
```

CONFIG\_ARM\_PSCI is not set ==> static inline bool psci\_smp\_available(void) { return false; }

==> 所以运行的是

```
smp_set_ops(mdsc->smp);
```

---

根据dts中cpus device\_node中的setting来初始化smp variables

in arch/arm/kernel/setup.c

```
setup_arch() --> arm_dt_init_cpu_maps()
```

```

1.  /*
2.   * arm_dt_init_cpu_maps - Function retrieves cpu nodes from the device tree
3.   * and builds the cpu logical map array containing MPIDR values related to
4.   * logical cpus
5.   *
6.   * Updates the cpu possible mask with the number of parsed cpu nodes
7.   */
8.  void __init arm_dt_init_cpu_maps(void)
9.  {
10.     /*
11.      * Temp logical map is initialized with UINT_MAX values that are
12.      * considered invalid logical map entries since the logical map must
13.      * contain a list of MPIDR[23:0] values where MPIDR[31:24] must
14.      * read as 0.
15.      */
16.     struct device_node *cpu, *cpus;
17.     int found_method = 0;
18.     u32 i, j, cpuidx = 1;
19.     u32 mpidr = is_smp() ? read_cpuid_mpidr() & MPIDR_HWID_BITMASK : 0; ①
20.
21.     u32 tmp_map[NR_CPUS] = { [0 ... NR_CPUS-1] = MPIDR_INVALID }; ②
22.     bool bootcpu_valid = false;
23.     cpus = of_find_node_by_path("/cpus");
24.                                     ③
25.
26.     if (!cpus)
27.         return;
28.
29.     for_each_child_of_node(cpus, cpu) {
30.                                     ④
31.         u32 hwid;
32.
33.         if (of_node_cmp(cpu->type, "cpu"))
34.                                     ⑤
35.             continue;
36.
37.         pr_debug(" * %s...\n", cpu->full_name);
38.         /*
39.          * A device tree containing CPU nodes with missing "reg"
40.          * properties is considered invalid to build the
41.          * cpu_logical_map.
42.          */
43.         if (of_property_read_u32(cpu, "reg", &hwid)) {
44.                                     ⑥
45.             pr_debug(" * %s missing reg property\n",
46.                     cpu->full_name);
47.             return;
48.         }
49.
50.         /*
51.          * 8 MSBs must be set to 0 in the DT since the reg property
52.          * defines the MPIDR[23:0].
53.          */

```

```

50.         if (hwid & ~MPIDR_HWID_BITMASK)
51.             return;
52.
53.         /*
54.          * Duplicate MPIDRs are a recipe for disaster.
55.          * Scan all initialized entries and check for
56.          * duplicates. If any is found just bail out.
57.          * temp values were initialized to UINT_MAX
58.          * to avoid matching valid MPIDR[23:0] values.
59.          */
60.         for (j = 0; j < cpuidx; j++)
61.             if (WARN(tmp_map[j] == hwid, "Duplicate /cpu reg "
62.                     "properties in the DT\n"))
63.                 return;
64.
65.         /*
66.          * Build a stashed array of MPIDR values. Numbering scheme
67.          * requires that if detected the boot CPU must be assigned
68.          * logical id 0. Other CPUs get sequential indexes starting
69.          * from 1. If a CPU node with a reg property matching the
70.          * boot CPU MPIDR is detected, this is recorded so that the
71.          * logical map built from DT is validated and can be used
72.          * to override the map created in smp_setup_processor_id().
73.          */
74.         if (hwid == mpidr) {
75.             i = 0;
76.             bootcpu_valid = true;
77.         } else {
78.             i = cpuidx++;
79.         }
80.
81.         if (WARN(cpuidx > nr_cpu_ids, "DT /cpu %u nodes greater than "
82.                 "max cores %u, capping them\n",
83.                 cpuidx, nr_cpu_ids)) {
84.             cpuidx = nr_cpu_ids;
85.             break;
86.         }
87.
88.         tmp_map[i] = hwid;
89.
90.         if (!found_method)
91.             found_method = set_smp_ops_by_method(cpu);
92.     }
93.
94.     /*
95.      * Fallback to an enable-method in the cpus node if nothing found in
96.      * a cpu node.
97.      */
98.     if (!found_method)
99.         set_smp_ops_by_method(cpus);
100.

```

```

101.         if (!bootcpu_valid) {
102.             pr_warn("DT missing boot CPU MPIDR[23:0], fall back to default cp
u_logical_map\n");
103.             return;
104.         }
105.
106.         /*
107.          * Since the boot CPU node contains proper data, and all nodes have
108.          * a reg property, the DT CPU list can be considered valid and the
109.          * logical map created in smp_setup_processor_id() can be overridden
110.          */
111.         for (i = 0; i < cpuidx; i++) {
112.             set_cpu_possible(i, true);
113.             cpu_logical_map(i) = tmp_map[i];
114.             pr_debug("cpu logical map 0x%x\n", cpu_logical_map(i));
115.         }
116.     }

```

①

read cp15's c0's Multiprocessor Affinity Register

```

1.  #define read_cpuid(reg) \
2.      ({ \
3.          unsigned int __val; \
4.          asm("mrc      p15, 0, %0, c0, c0, " __stringify(reg) \
5.              : "=r" (__val) \
6.              : \
7.              : "cc"); \
8.          __val; \
9.      })
10.
11. static inline unsigned int __attribute_const__ read_cpuid_mpidr(void)
12. {
13.     return read_cpuid(CPUID_MPIDR);
14. }

```

==>

mrc p15, 0, Rx, c0, c0, 5

mpidr = Rx

mpidr的byte 0 = Affinity Level 0, byte 1 = Affinity Level 1, byte 2 = Affinity Level 2, byte 3必须为0

所以这里的mask MPIDR\_HWID\_BITMASK = 0x00FFFFFF

在gemstone2上 mpidr = 0x80ffff00

bit 31 = 1, Indicates that the processor implements the Multiprocessing Extensions register format.

②

tmp\_map[NR\_CPUS] = tmp\_map[2] = MPIDR\_INVALID = 0xFF000000

③

get cpus device\_node as follow

```
1.      cpus {
2.          #address-cells = <0x1>;
3.          #size-cells = <0x0>;
4.
5.          cpu@0 {
6.              device_type = "cpu";
7.              compatible = "arm,cortex-a53";
8.              reg = <0xffff00>;
9.          };
10.
11.         cpu@1 {
12.             device_type = "cpu";
13.             compatible = "arm,cortex-a53";
14.             reg = <0xffff01>;
15.             enable-method = "marvell,pegmatite-apmu-boot";
16.         };
17.     };
```

④

比如gemstone2, cpus device\_node有2个childs。对child进行enumeration

⑤

对child的合法性检查, child必须有

```
device_type = "cpu";
```

这个property。

⑥

read "reg" property from child, hwid = 0xffff00 或0xffff01

⑦

这里是检查各个child中reg property的值不能有重复的。

⑧

core cpu@0 (reg = <0xffff00>)是boot core,也就是当前运行的core,core cpu@1还没有enable呢

⑨

对core cpu@1而言

```

1. extern struct of_cpu_method __cpu_method_of_table[];
2.
3. static const struct of_cpu_method __cpu_method_of_table_sentinel
4.     __used __section(__cpu_method_of_table_end);
5.
6. static int __init set_smp_ops_by_method(struct device_node *node)
7. {
8.     const char *method;
9.     struct of_cpu_method *m = __cpu_method_of_table;
10.
11.     if (of_property_read_string(node, "enable-method", &method))
12.         return 0;
13.
14.     for (; m->method; m++)
15.         if (!strcmp(m->method, method)) {
16.             smp_set_ops(m->ops);
17.             return 1;
18.         }
19.
20.     return 0;
21. }

```

查询\_\_cpu\_method\_of\_table[].

in drivers/platform/pegmatite/smp/platsmp.c

```

1. struct smp_operations pegmatite_smp_ops __initdata = {
2.     .smp_prepare_cpus      = pegmatite_smp_prepare_cpus,
3.     .smp_boot_secondary   = pegmatite_boot_secondary,
4. #ifdef CONFIG_HOTPLUG_CPU
5.     .cpu_die               = pegmatite_cpu_die,
6.     .cpu_kill              = pegmatite_cpu_kill,
7. #endif
8. };
9.
10. ....
11.
12. CPU_METHOD_OF_DECLARE(pegmatite_smp, "marvell,pegmatite-apmu-boot", &pegmatite_smp_ops);

```

set\_smp\_ops\_by\_method()中的



```
smp_set_ops(m->ops);
```

```
m->ops = pegmatite_smp_ops
```

==>

in arch/arm/kernel/smp.c

```
1.  static struct smp_operations smp_ops;  
2.  
3.  void __init smp_set_ops(struct smp_operations *ops)  
4.  {  
5.      if (ops)  
6.          smp_ops = *ops;  
7.  };
```

使得global variable smp\_ops = pegmatite\_smp\_ops

初始化smp\_ops , 在smp initialization II会用到它。

⑩

对gemstone2而言 , cpu mask的bit 0, bit 1置位 , 表示有2 cores.

```
cpu_logical_map(0) = 0xffff00
```

```
cpu_logical_map(1) = 0xffff01
```