

in kernel/smp.c

```
1.  /* Called by boot processor to activate the rest. */
2.  void __init smp_init(void)
3.  {
4.      unsigned int cpu;
5.
6.      idle_threads_init(); ①
7.
8.      /* FIXME: This should be done in userspace --RR */
9.      for_each_present_cpu(cpu) {
10.         if (num_online_cpus() >= setup_max_cpus)
11.             break;
12.         if (!cpu_online(cpu))
13.             cpu_up(cpu); ②
14.     }
15.
16.     /* Any cleanup work */
17.     smp_announce();
18.     smp_cpus_done(setup_max_cpus);
19. }
```

①

Linux kernel为每个core维护着一个idle process，即如果没有任何task可运行时，就运行它。

所以这里要为即将满血复活的secondary core生成对应的idle process。

in kernel/smpboot.c

```
1.  /**
2.   * idle_threads_init - Initialize idle threads for all cpus
3.   */
4.  void __init idle_threads_init(void)
5.  {
6.      unsigned int cpu, boot_cpu;
7.
8.      boot_cpu = smp_processor_id();
9.
10.     for_each_possible_cpu(cpu) {
11.         if (cpu != boot_cpu)           (A)
12.             idle_init(cpu);           (B)
13.     }
14. }
```

(A)

排除boot cpu

(B)

```

1.  /**
2.   * idle_init - Initialize the idle thread for a cpu
3.   * @cpu:      The cpu for which the idle thread should be initialized
4.   *
5.   * Creates the thread if it does not exist.
6.   */
7.  static inline void idle_init(unsigned int cpu)
8.  {
9.      struct task_struct *tsk = per_cpu(idle_threads, cpu);
10.         (C)
11.
12.      if (!tsk) {
13.          tsk = fork_idle(cpu);
14.          (D)
15.          if (IS_ERR(tsk))
16.              pr_err("SMP: fork_idle() failed for CPU %u\n", cpu);
17.          else
18.              per_cpu(idle_threads, cpu) = tsk;
19.      }
20.  }
21.  struct task_struct *fork_idle(int cpu)
22.  {
23.      struct task_struct *task;
24.      task = copy_process(CLONE_VM, 0, 0, NULL, &init_struct_pid, 0);
25.      (E)
26.      if (!IS_ERR(task)) {
27.          init_idle_pids(task->pids);
28.          init_idle(task, cpu);
29.      }
30.      return task;
31.  }

```

(C)

idle process的task_struct *是个per_cpu variable。

(D)

create idle process

(E)

idle process完全运行在kernel mode，所以它与boot cpu的idle process是share virtual

memory的。

当secondary core初始化好以后，运行的第一个process code就是该core对应的idle process。

②

in kernel/cpu.c

```

1. int cpu_up(unsigned int cpu)
2. {
3.     int err = 0;
4.
5.     if (!cpu_possible(cpu)) {
6.         pr_err("can't online cpu %d because it is not configured as may-h
otadd at boot time\n",
7.             cpu);
8. #if defined(CONFIG_IA64)
9.         pr_err("please check additional_cpus= boot parameter\n");
10. #endif
11.         return -EINVAL;
12.     }
13.
14.     err = try_online_node(cpu_to_node(cpu));
15.     if (err)
16.         return err;
17.
18.     cpu_maps_update_begin();
19.
20.     if (cpu_hotplug_disabled) {
21.         err = -EBUSY;
22.         goto out;
23.     }
24.
25.     err = _cpu_up(cpu, 0);
26.
27. out:
28.     cpu_maps_update_done();
29.     return err;
30. }
31.
32. static int _cpu_up(unsigned int cpu, int tasks_frozen)
33. {
34.     int ret, nr_calls = 0;
35.     void *hcpu = (void *) (long)cpu;
36.     unsigned long mod = tasks_frozen ? CPU_TASKS_FROZEN : 0;
37.     struct task_struct *idle;
38.
39.     cpu_hotplug_begin();
40.
41.     if (cpu_online(cpu) || !cpu_present(cpu)) {
42.         ret = -EINVAL;
43.         goto out;
44.     }
45.
46.     idle = idle_thread_get(cpu);
47.
48.     if (IS_ERR(idle)) {
49.         ret = PTR_ERR(idle);
50.         goto out;
51.     }

```

(A)

(B)

```

51.
52.     ret = smpboot_create_threads(cpu);
53.                                     (C)
54.     if (ret)
55.         goto out;
56.     ret = __cpu_notify(CPU_UP_PREPARE | mod, hcpu, -1, &nr_calls);
57.                                     (D)
58.     if (ret) {
59.         nr_calls--;
60.         pr_warn("%s: attempt to bring up CPU %u failed\n",
61.                 __func__, cpu);
62.         goto out_notify;
63.     }
64.     /* Arch-specific enabling code. */
65.     ret = __cpu_up(cpu, idle);
66.                                     (E)
67.     if (ret != 0)
68.         goto out_notify;
69.     BUG_ON(!cpu_online(cpu));
70.     /* Wake the per cpu threads */
71.     smpboot_unpark_threads(cpu);
72.                                     (F)
73.     /* Now call notifier in preparation. */
74.     cpu_notify(CPU_ONLINE | mod, hcpu);
75.                                     (G)
76. out_notify:
77.     if (ret != 0)
78.         __cpu_notify(CPU_UP_CANCELED | mod, hcpu, nr_calls, NULL);
79. out:
80.     cpu_hotplug_done();
81.
82.     return ret;
83. }

```

(A)

如果该core已经活了或不存在，自然skip。

(B)

取得对应core的idle process指针(task_struct *),也就是前一步创建的idle process。

(C)

每个core还有hotplug thread , 没研究。

(D)

observer pattern的应用。发出CPU_UP_PREPARE notification。

(E)

这是boot secondary core的核心,见后分析。如果该function成功返回,那么该core就online了。

(F)

同样是cpu hotplug相关操作

(G)

observer pattern的应用。发出CPU_ONLINE notification。